

# CASCADE

A Just-in-Time Compiler for FPGAs

ASPLOS 2019

# Opportunities and Challenges

## Performance:

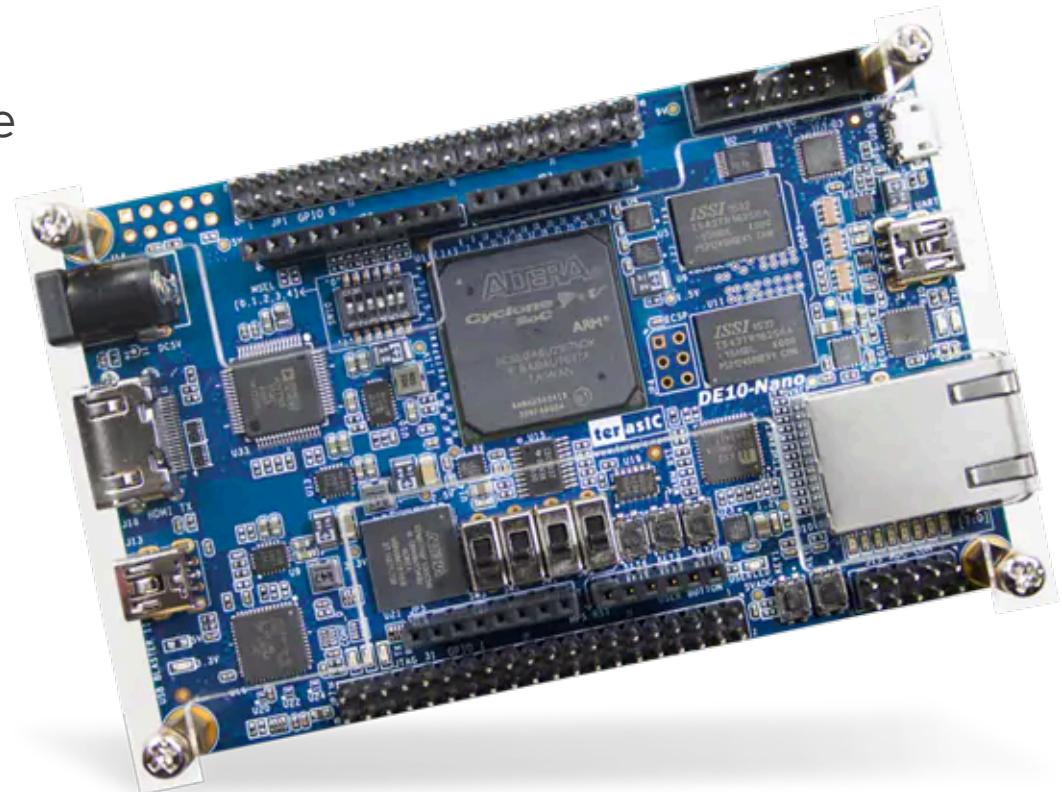
- Faster than CPUs by several orders of magnitude
- Lower cost and time to market than ASICs

## Availability:

- Increasingly common in the data-center

## Difficult to Program:

- Compilation can take hours or longer!
- Expensive, baroque IDEs!
- No open source alternatives!



# Just-in-Time Compilation

## 1. Run Code in a Simulator:

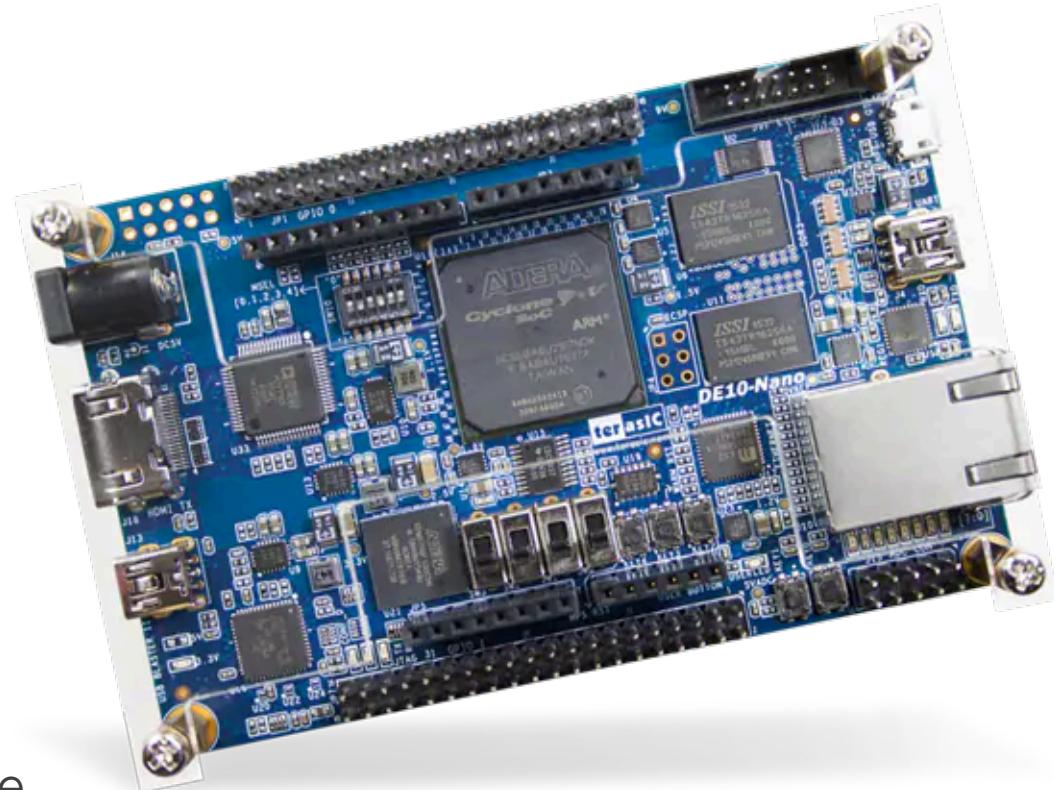
- No waiting for the compiler
- Shortened compile-test-debug cycle

## 2. Compile in the Background:

- Treat closed-source compilers as a black box
- Replace the IDE with a light-weight front-end

## 3. Transition to Hardware When Finished:

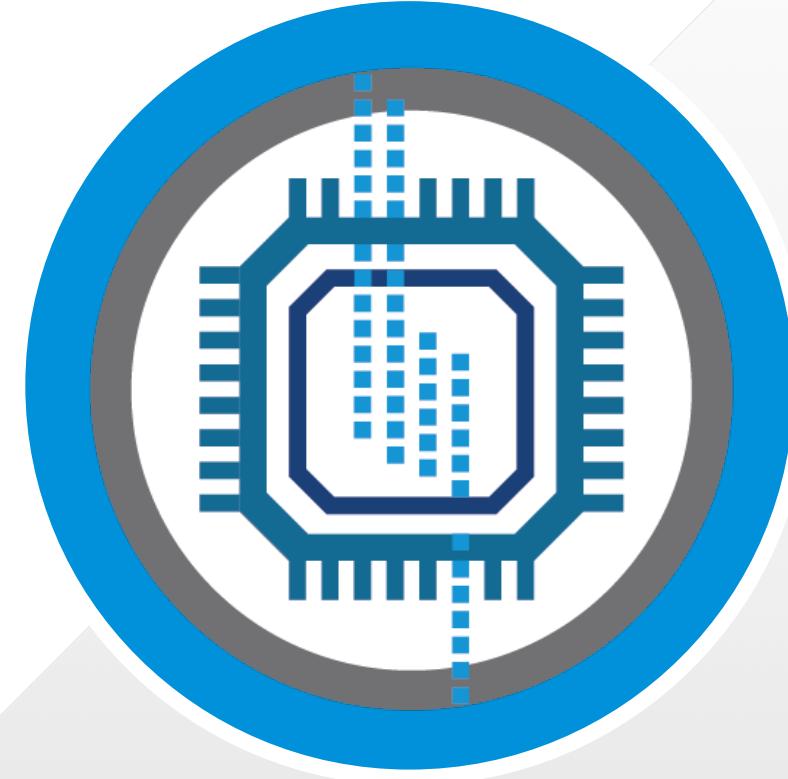
- Runtime manages low-level details
- Code gets faster over time as it reaches hardware



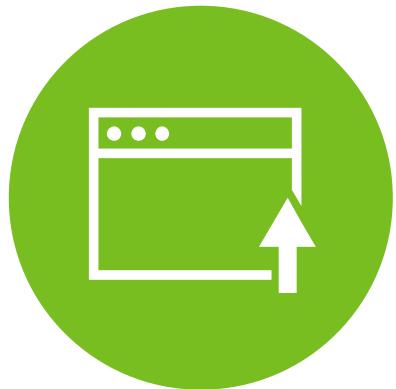
# CASCADE:

The First Just-in-Time Compiler for FPGAs

<https://github.com/vmware/cascade>

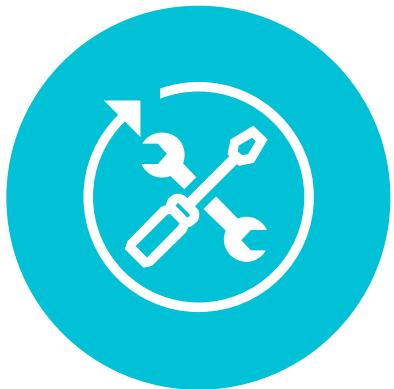


# Design Goals



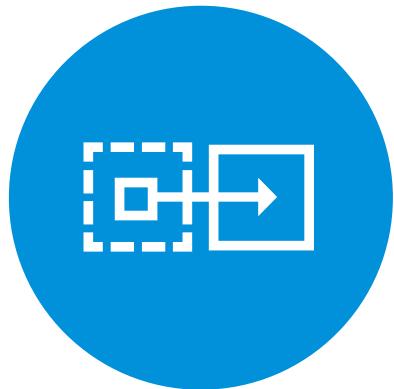
## Interactivity

Run code  
immediately as it's  
written.



## Expressiveness

Printf-style  
debugging from  
hardware.



## Portability

Write code once.  
Run it on many  
platforms.



## Performance

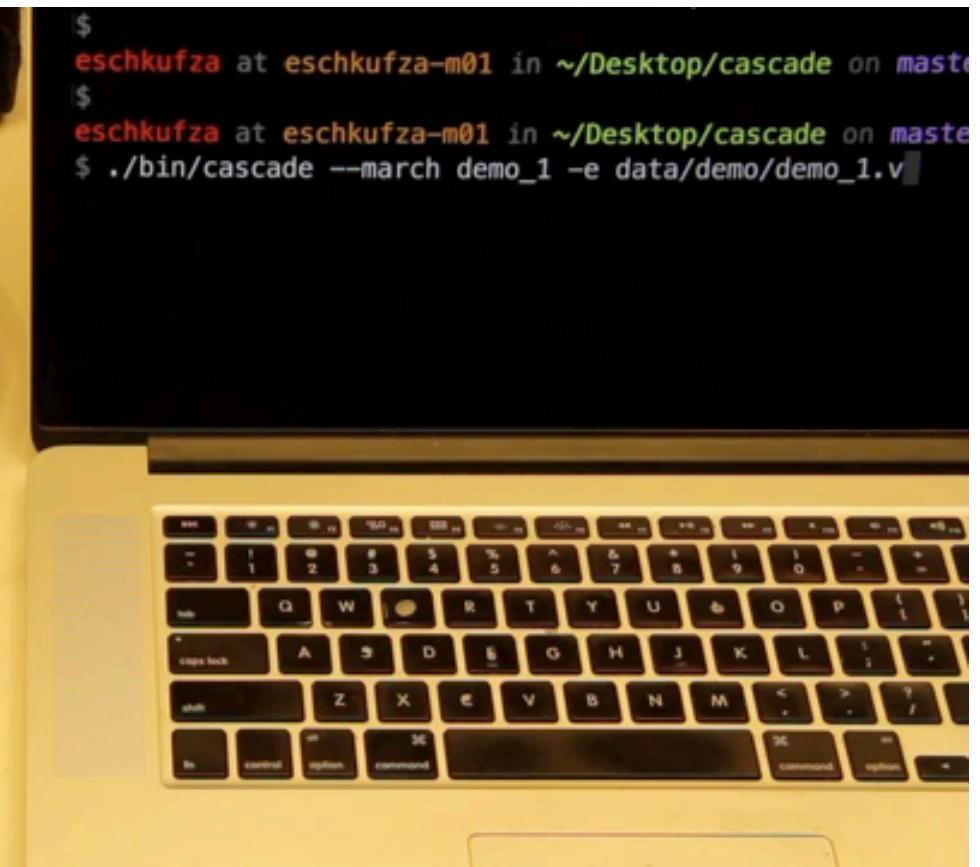
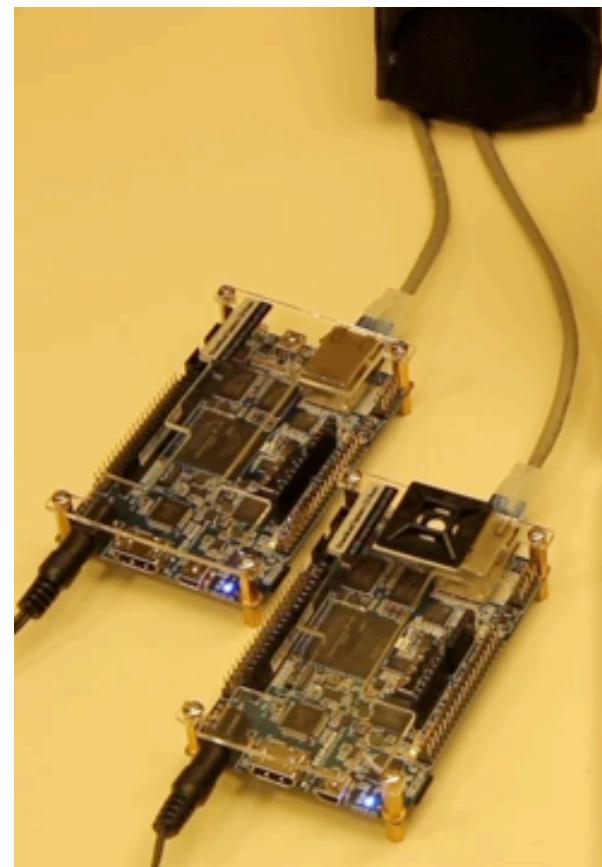
Don't pay for what  
you don't use.

# Basic User Interaction

Code is compiled continuously.

Side effects are immediately visible in hardware.

```
Macintosh HD — top — 80x24  
assign led.val = pad.val;
```



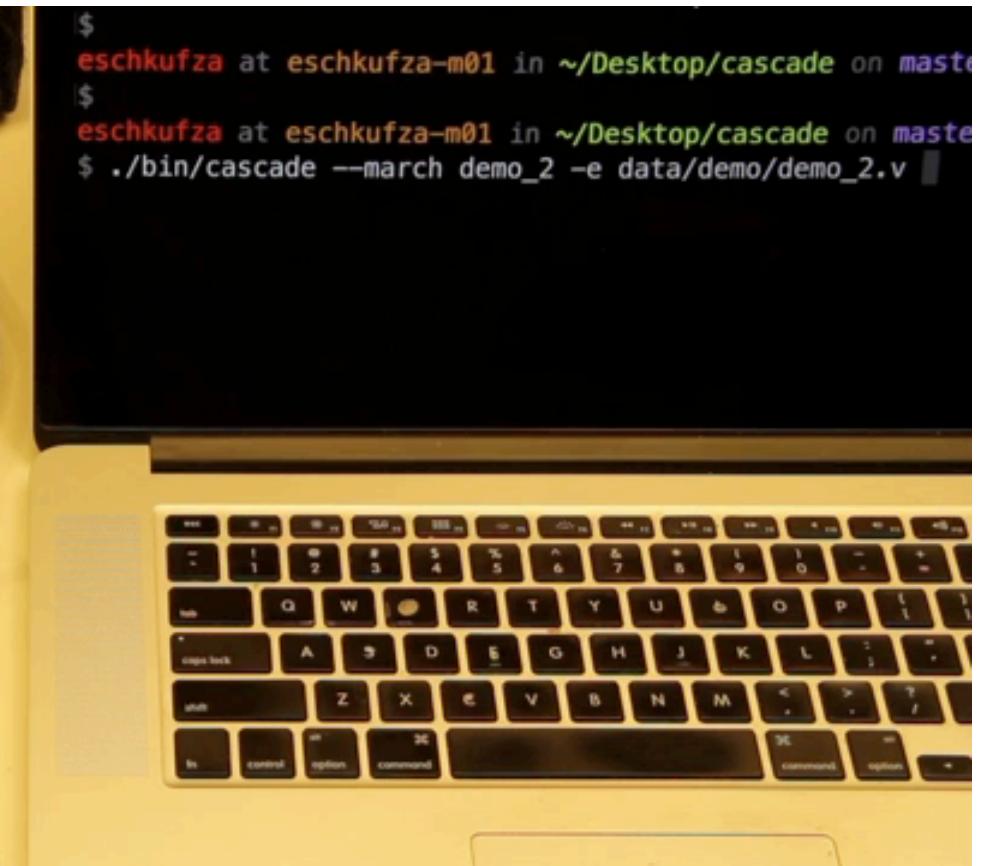
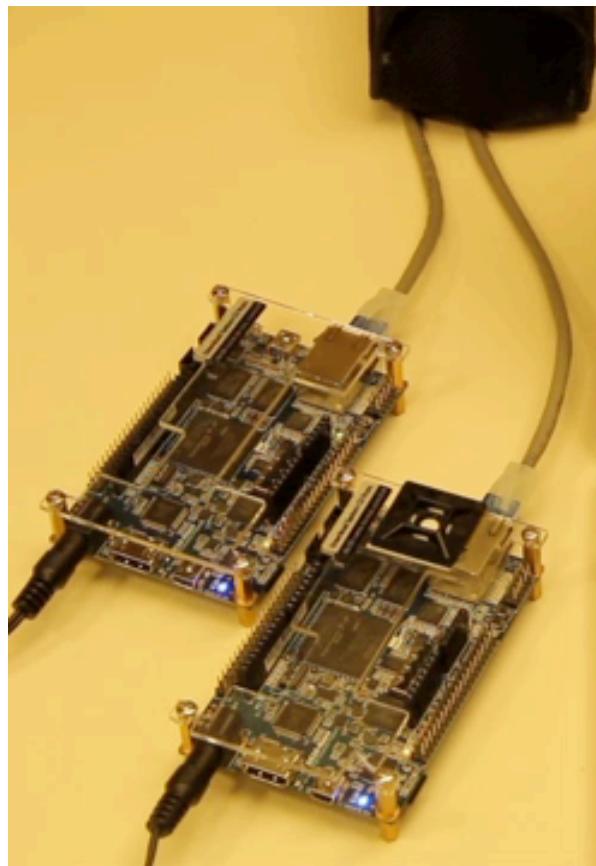
# Hardware printf

A message is printed to the terminal when the user interacts with the buttons.

```
Macintosh HD — top — 80x24

assign led.val = pad.val;

reg r = 0;
always @(posedge pad.val) begin
  r <= !r;
  $display(r ? "tick" : "tock");
end
```



# Why Cascade Works

```
1: procedure EVAL(e)
2:   if e is an update then
3:     perform sequential update
4:   else
5:     evaluate combinational logic
6:   end if
7:   enqueue new events
8: end procedure
```

```
1: procedure REFERENCESCHEDULER
2:   while  $\top$  do
3:     if  $\exists$  activated events then
4:       EVAL(any activated event)
5:     else if  $\exists$  update events then
6:       activate all update events
7:     else
8:       advance time t; schedule recurring events
9:     end if
10:   end while
11: end procedure
```

## Simulation Reference Model:

- Verilog defined in terms of scheduler algorithm
- Any system which produces the same set of side effects is a well-defined model for Verilog!

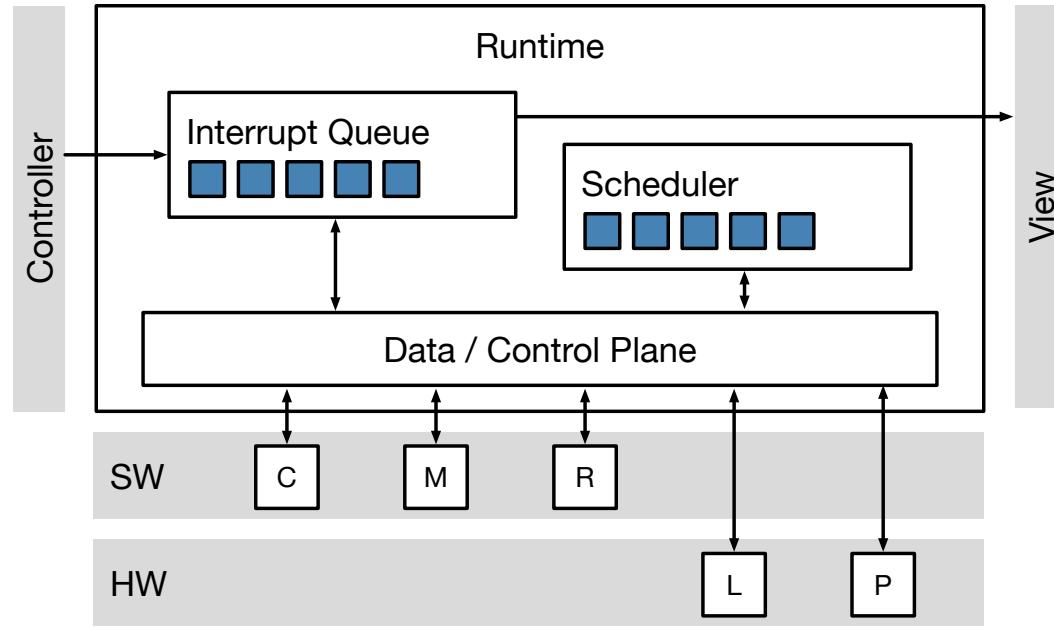
## Small-Step Semantics:

- Events processed one at a time
- Observable sequence points give illusion that they take place simultaneously

## Safe Windows:

- Well-defined places where state is stable
- Soundly swap implementations, eval new code, emit print statements, etc...

# Runtime Architecture



## Distributed-System-Like IR:

- Composed of Verilog sub-programs
- Communication over a constrained protocol

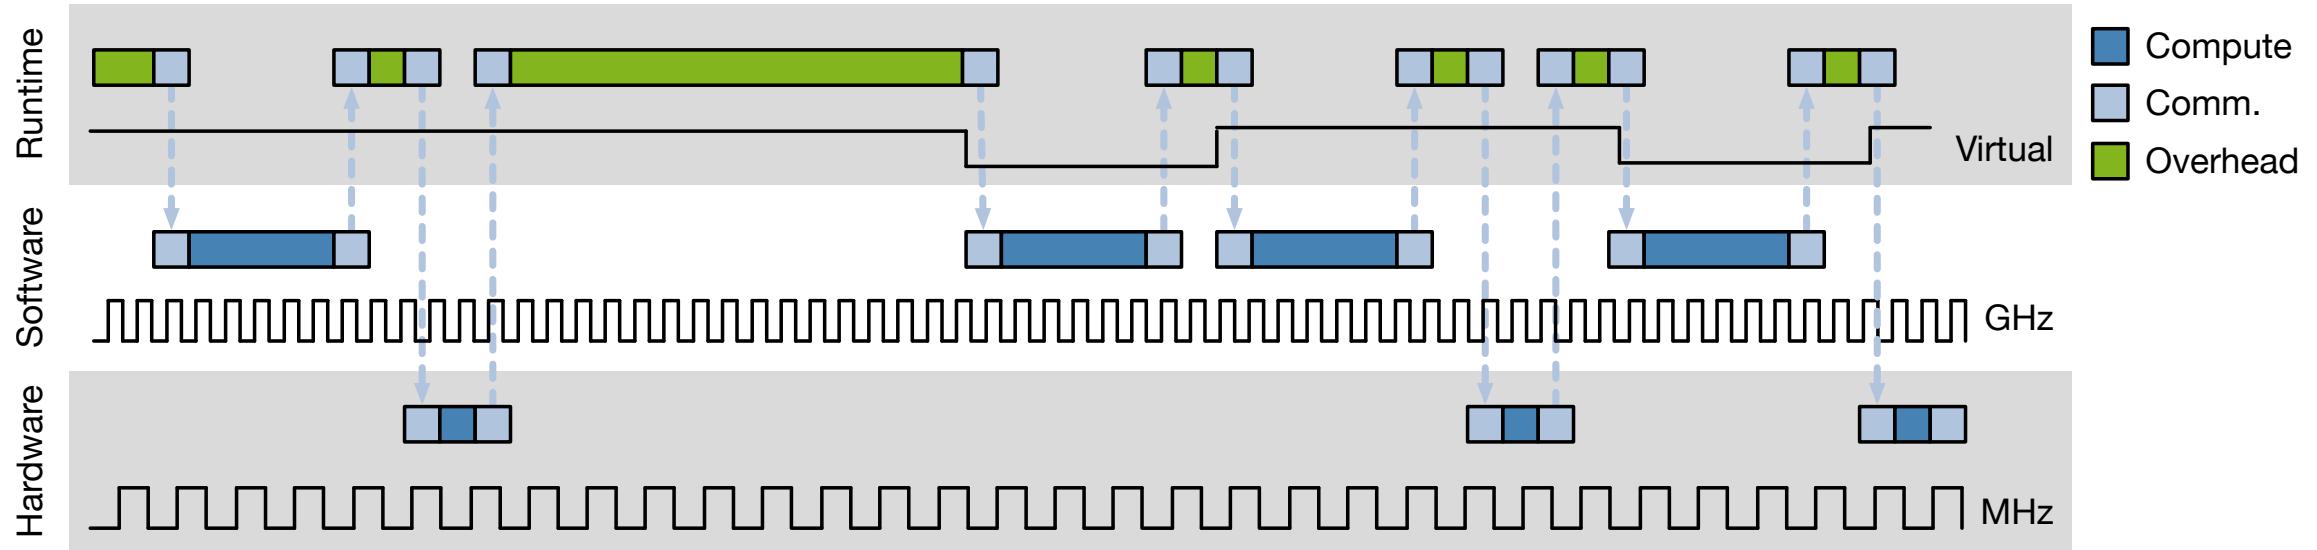
## Monadic Semantics:

- Allows dynamic swapping of implementations
- Agnostic to whether modules are located in hardware or software

## Data/Control Plane:

- Module state stored in implementation-specific “engines”
- Communication via message passing interface

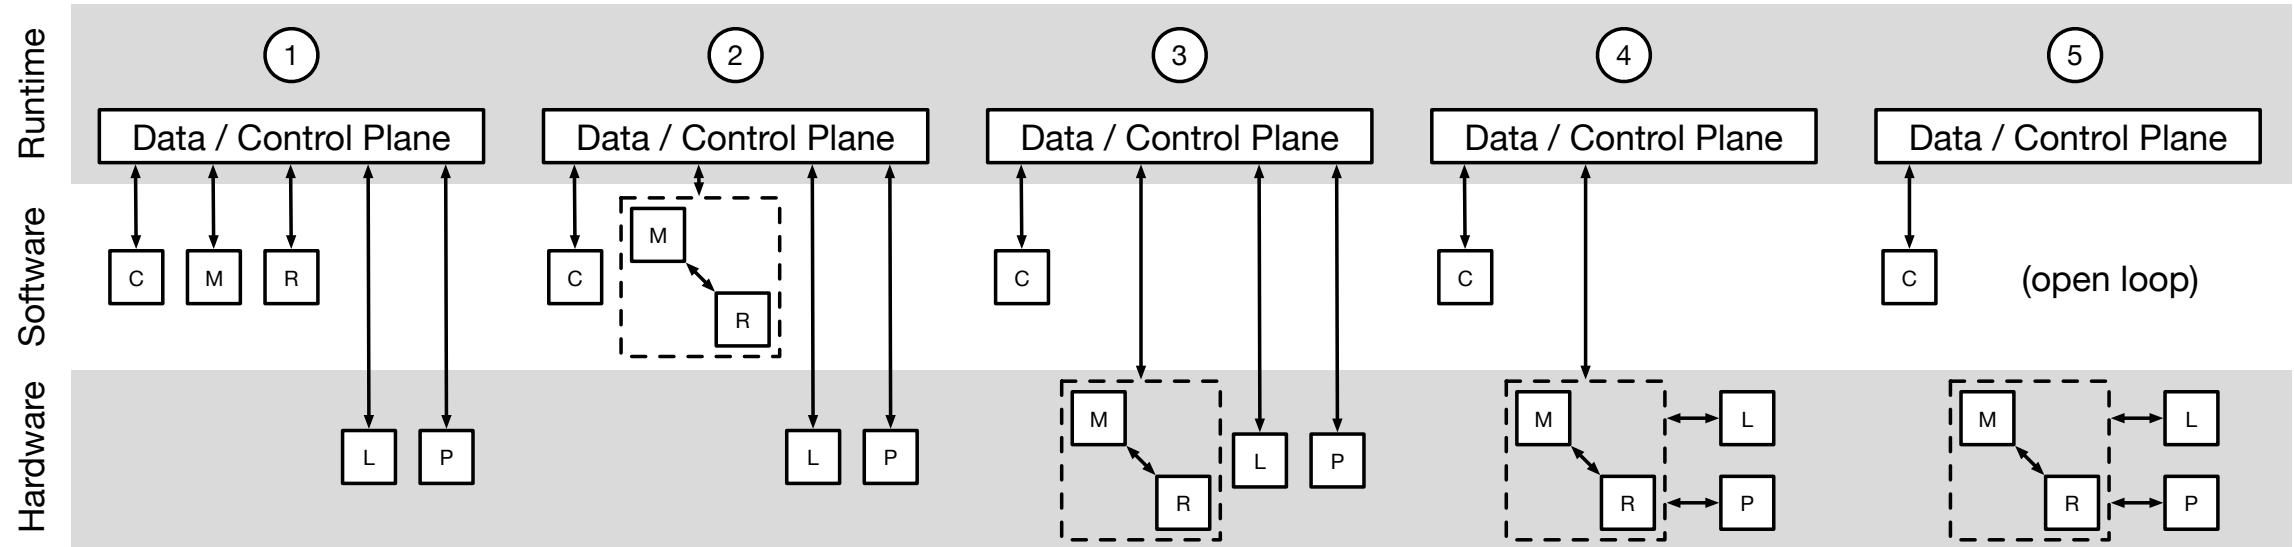
# Performance Contract



## Multiple Clock Domains:

- Compute at different frequencies, losses due to overhead
- Maximize the number of scheduler iterations per second (virtual clock frequency)

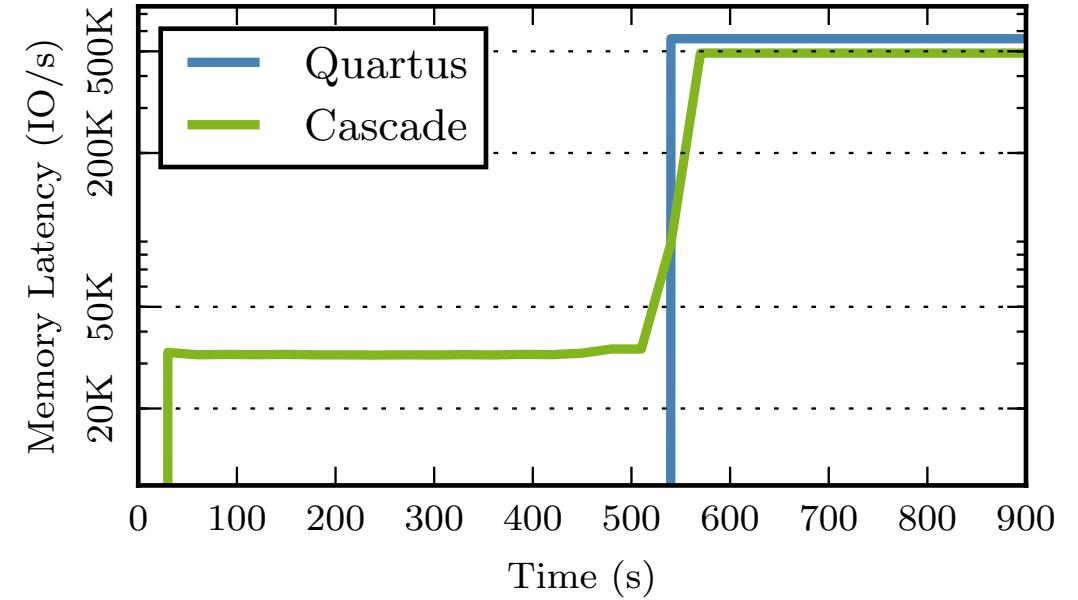
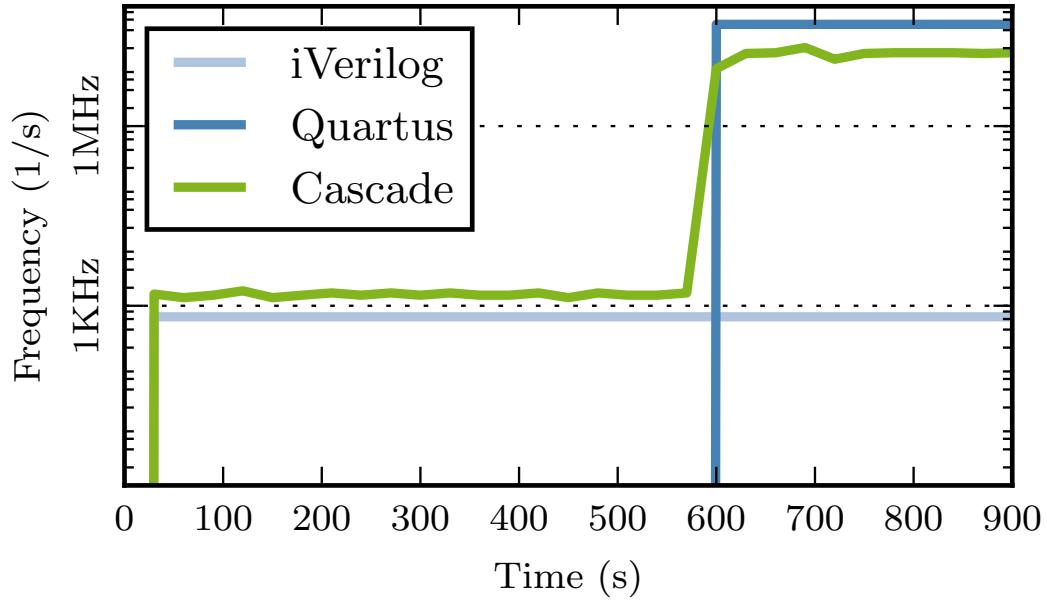
# Optimization Pipeline



## High-Level Concepts:

- Gradual transition from software to hardware
- Dynamic inlining removes overhead of runtime-mediated communication
- Open-loop execution removes overhead of runtime-mediated scheduling

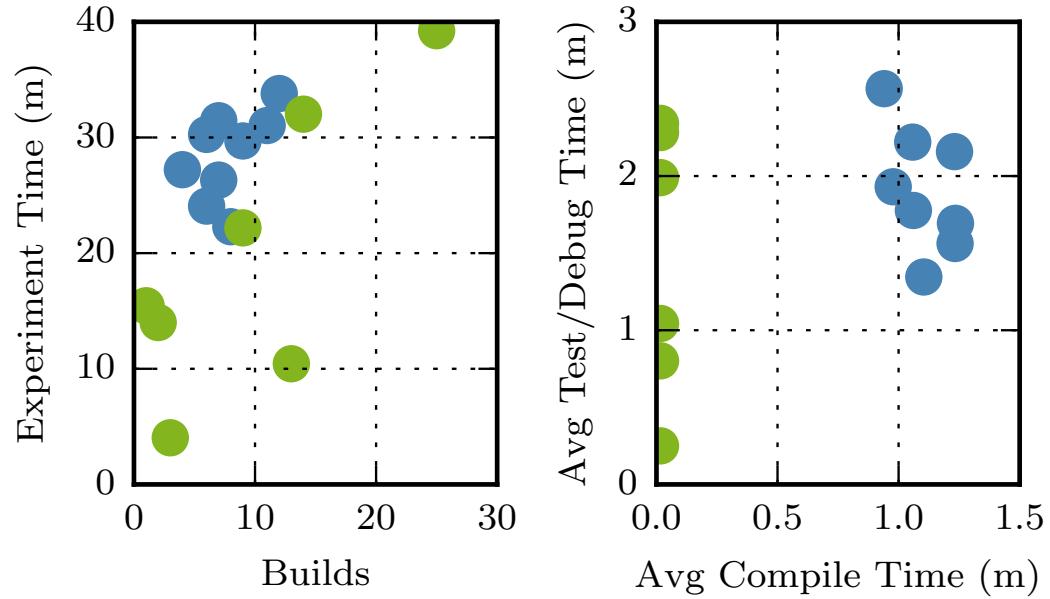
# Experiments



## Benchmarks:

- Compute bound (bitcoin miner) and communication bound (regular expression streaming)
- Code begins execution with performance comparable to simulator, approaches native performance in the limit, retains support for unsynthesizable Verilog

# User Study



## Debugging Task:

- Users given spec and buggy code
- Asked to produce a working solution

## Sample Set:

- 20 participants ranging from novice to expert
- Half given Cascade, half given Quartus

## Observations:

- Users performed more compilation cycles and finished task faster when using Cascade
- Spent same time testing and debugging
- Encourages more compilation, not sloppy coding



# Thank You

<https://github.com/vmware/cascade>

<https://research.vmware.com>