

# Stochastic Optimization of Floating-point Programs with Tunable Precision

eric schkufza, rahul sharma, alex aiken  
stanford university

# Overview

- **Floating-point kernels:** Nearly ubiquitous in high-performance computing
- **Getting the best performance:** Requires full exploitation of the dark corners of an instruction set and how it interacts with machine resources
- **But we can do even better:** Give up on full precision for applications that don't need it

# Overview

- **Stochastic Optimization:** Automated method for generating high-performance kernels that trade a reduction in precision for improvements in code size and runtime

# Goals and Related Work

- **Mixed fixed / floating-point kernels:** Build on previous work [schkufza 13, sharma 13]
- **Improved approximation:** Give user finer-grained control [lam 13, rubio-gonzalez 13]
- **Assembly level:** Direct optimization, no further trusted optimization kernel [sidiroglou-douskos 11]
- **Optimize kernels:** Improve end-to-end performance automatically [chilimbi 10, zou 12]

# Example

expert

```
vmovddup    %xmm0, %xmm0
vmulpd      (%rdi), %xmm0, %xmm2
vroundpd    $0, %xmm2, %xmm2
vmulpd      0x10(%rdi), %xmm2, %xmm1
vcvtpd2dqx %xmm2, %xmm3
vmulpd      0x20(%rdi), %xmm2, %xmm2
vaddpd      %xmm1, %xmm0, %xmm1
vmovapd     0x30(%rdi), %xmm0
vpaddb      0x40(%rdi), %xmm3, %xmm3
vpsllb      $20, %xmm3, %xmm3
vpshufb     $114, %xmm3, %xmm3
vaddpd      %xmm2, %xmm1, %xmm1
vmulpd      0x50(%rdi), %xmm1, %xmm2
vaddpd      0x60(%rdi), %xmm2, %xmm2
vmulpd      %xmm1, %xmm2, %xmm2
vaddpd      0x70(%rdi), %xmm2, %xmm2
vmulpd      %xmm1, %xmm2, %xmm2
vaddpd      0x80(%rdi), %xmm2, %xmm2
vmulpd      %xmm1, %xmm2, %xmm2
vaddpd      0x90(%rdi), %xmm2, %xmm2
vmulpd      %xmm1, %xmm2, %xmm2
vaddpd      0xa0(%rdi), %xmm2, %xmm2
vmulpd      %xmm1, %xmm2, %xmm2
vaddpd      0xb0(%rdi), %xmm2, %xmm2
vmulpd      %xmm1, %xmm2, %xmm2
vaddpd      0xc0(%rdi), %xmm2, %xmm2
vmulpd      %xmm1, %xmm2, %xmm2
vaddpd      0xd0(%rdi), %xmm2, %xmm2
vmulpd      %xmm1, %xmm2, %xmm2
vaddpd      0xe0(%rdi), %xmm2, %xmm2
vmulpd      %xmm1, %xmm2, %xmm2
vaddpd      0xf0(%rdi), %xmm2, %xmm2
vmulpd      %xmm1, %xmm2, %xmm2
vaddpd      %xmm0, %xmm2, %xmm2
vmulpd      %xmm1, %xmm2, %xmm1
vaddpd      %xmm0, %xmm1, %xmm0
vmulpd      %xmm3, %xmm0, %xmm0
retq
```

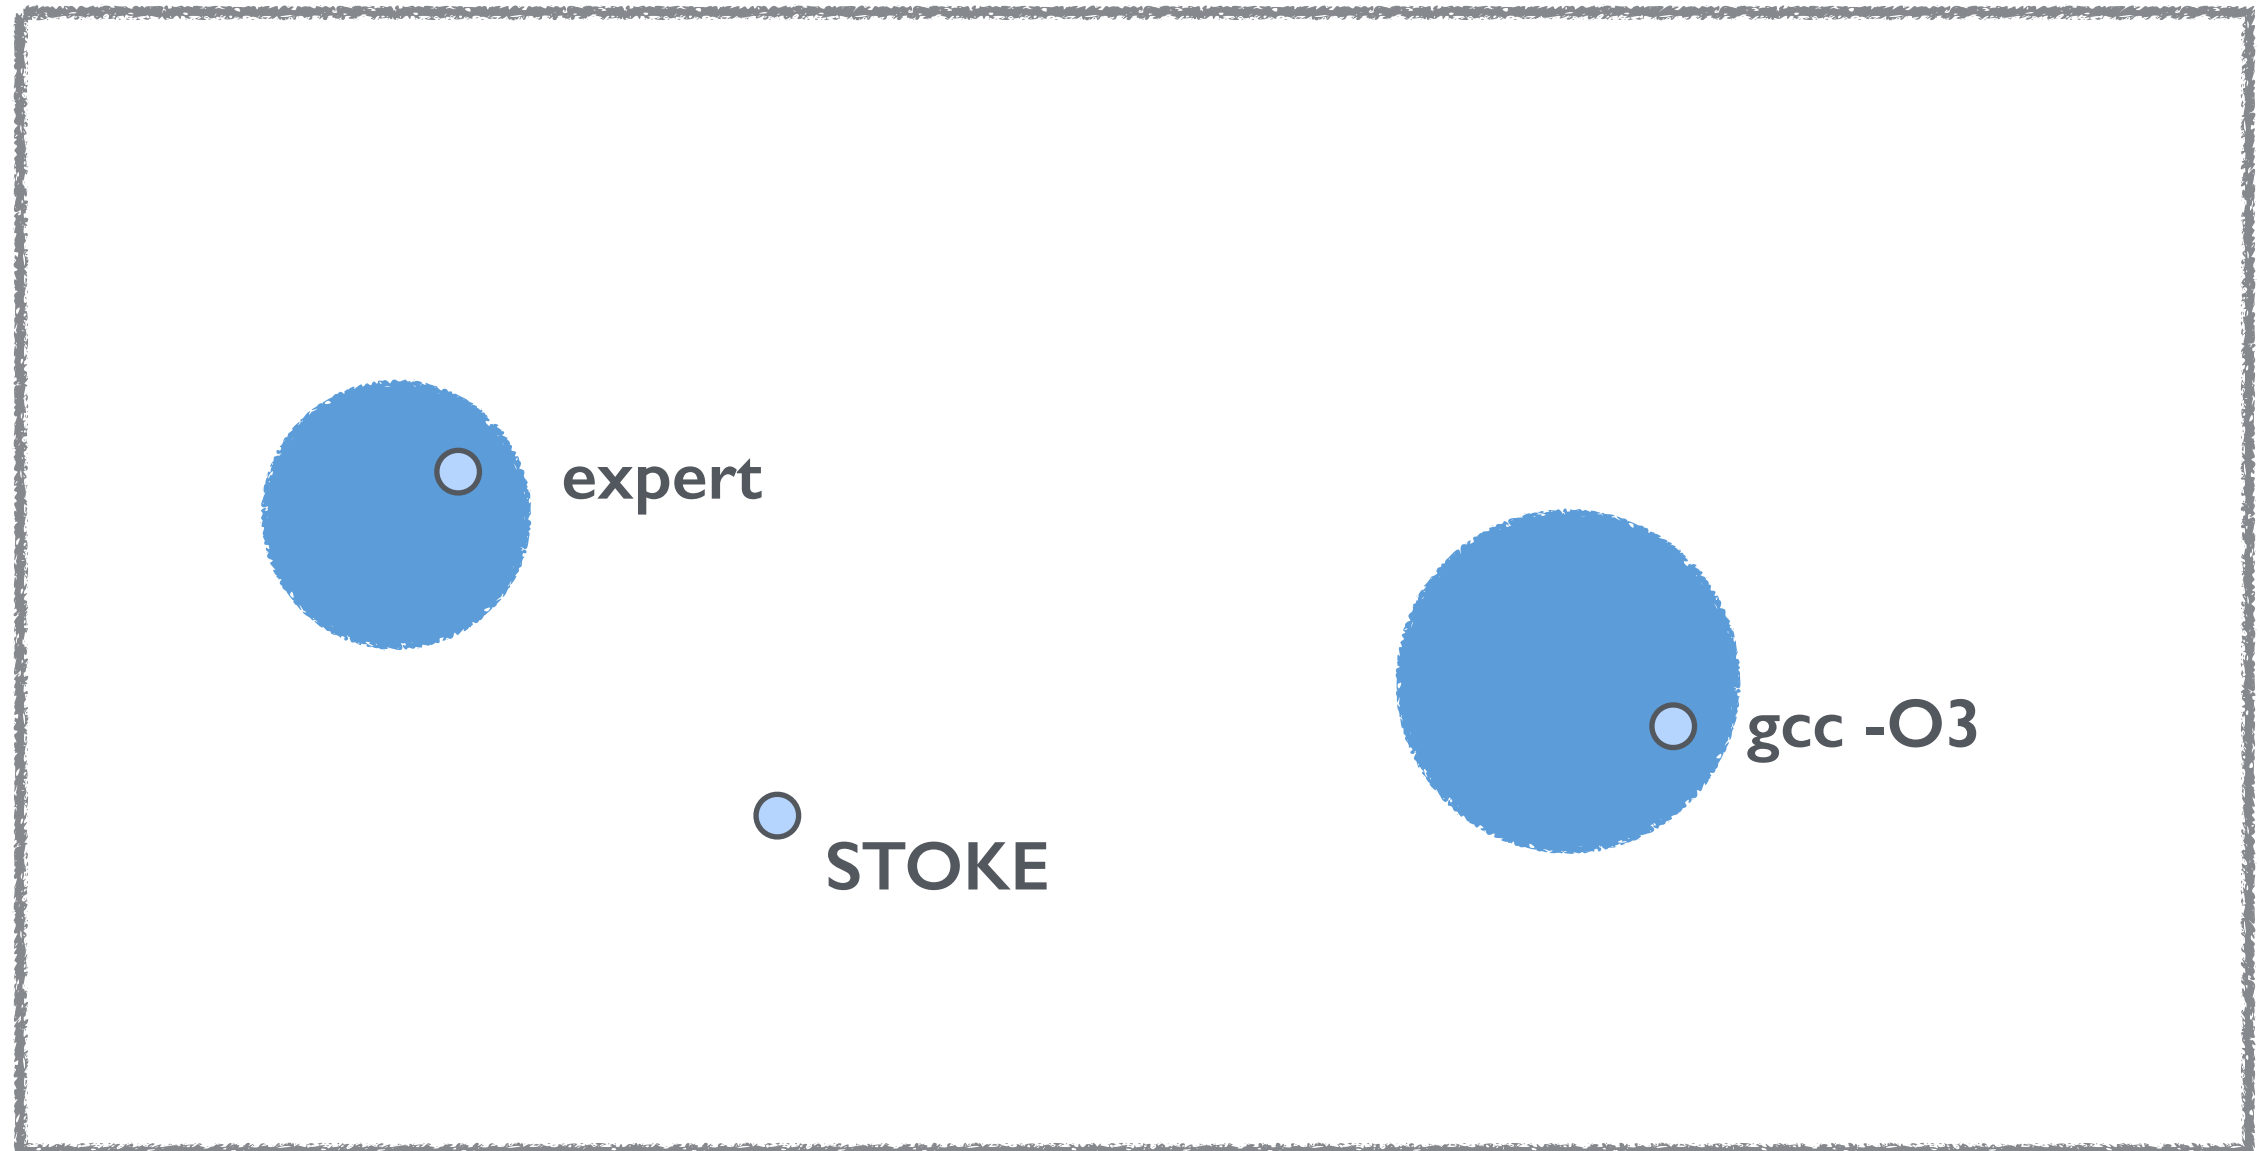
STOKE

```
vmulpd      (%rdi), %xmm0, %xmm2
vroundpd    $0xfffffffffffffffe, %xmm2, %xmm2
vcvttdpd2dq %xmm2, %xmm3
vmulpd      0x10(%rdi), %xmm2, %xmm1
vlddqu      0x90(%rdi), %xmm2
vaddpd      %xmm1, %xmm0, %xmm1
vmulpd      %xmm1, %xmm2, %xmm2
vpaddw      0x40(%rdi), %xmm3, %xmm3
vmovapd     0x30(%rdi), %xmm0
vpsllq      $0x14, %xmm3, %xmm3
vaddpd      0xa0(%rdi), %xmm2, %xmm2
vmulpd      %xmm1, %xmm2, %xmm2
vaddpd      0xb0(%rdi), %xmm2, %xmm2
vmulsd      %xmm1, %xmm2, %xmm2
vaddpd      0xc0(%rdi), %xmm2, %xmm2
vmulpd      %xmm1, %xmm2, %xmm2
vaddpd      0xd0(%rdi), %xmm2, %xmm2
vmulpd      %xmm1, %xmm2, %xmm2
vaddpd      0xe0(%rdi), %xmm2, %xmm2
vmulsd      %xmm1, %xmm2, %xmm2
vaddpd      0xf0(%rdi), %xmm2, %xmm2
vmulsd      %xmm1, %xmm2, %xmm2
vaddsd      %xmm0, %xmm2, %xmm2
vpshufd     $0x3, %xmm3, %xmm3
vmulpd      %xmm1, %xmm2, %xmm1
vaddsd      %xmm0, %xmm1, %xmm0
vmulpd      %xmm3, %xmm0, %xmm0
retq
```

# Optimization Notes

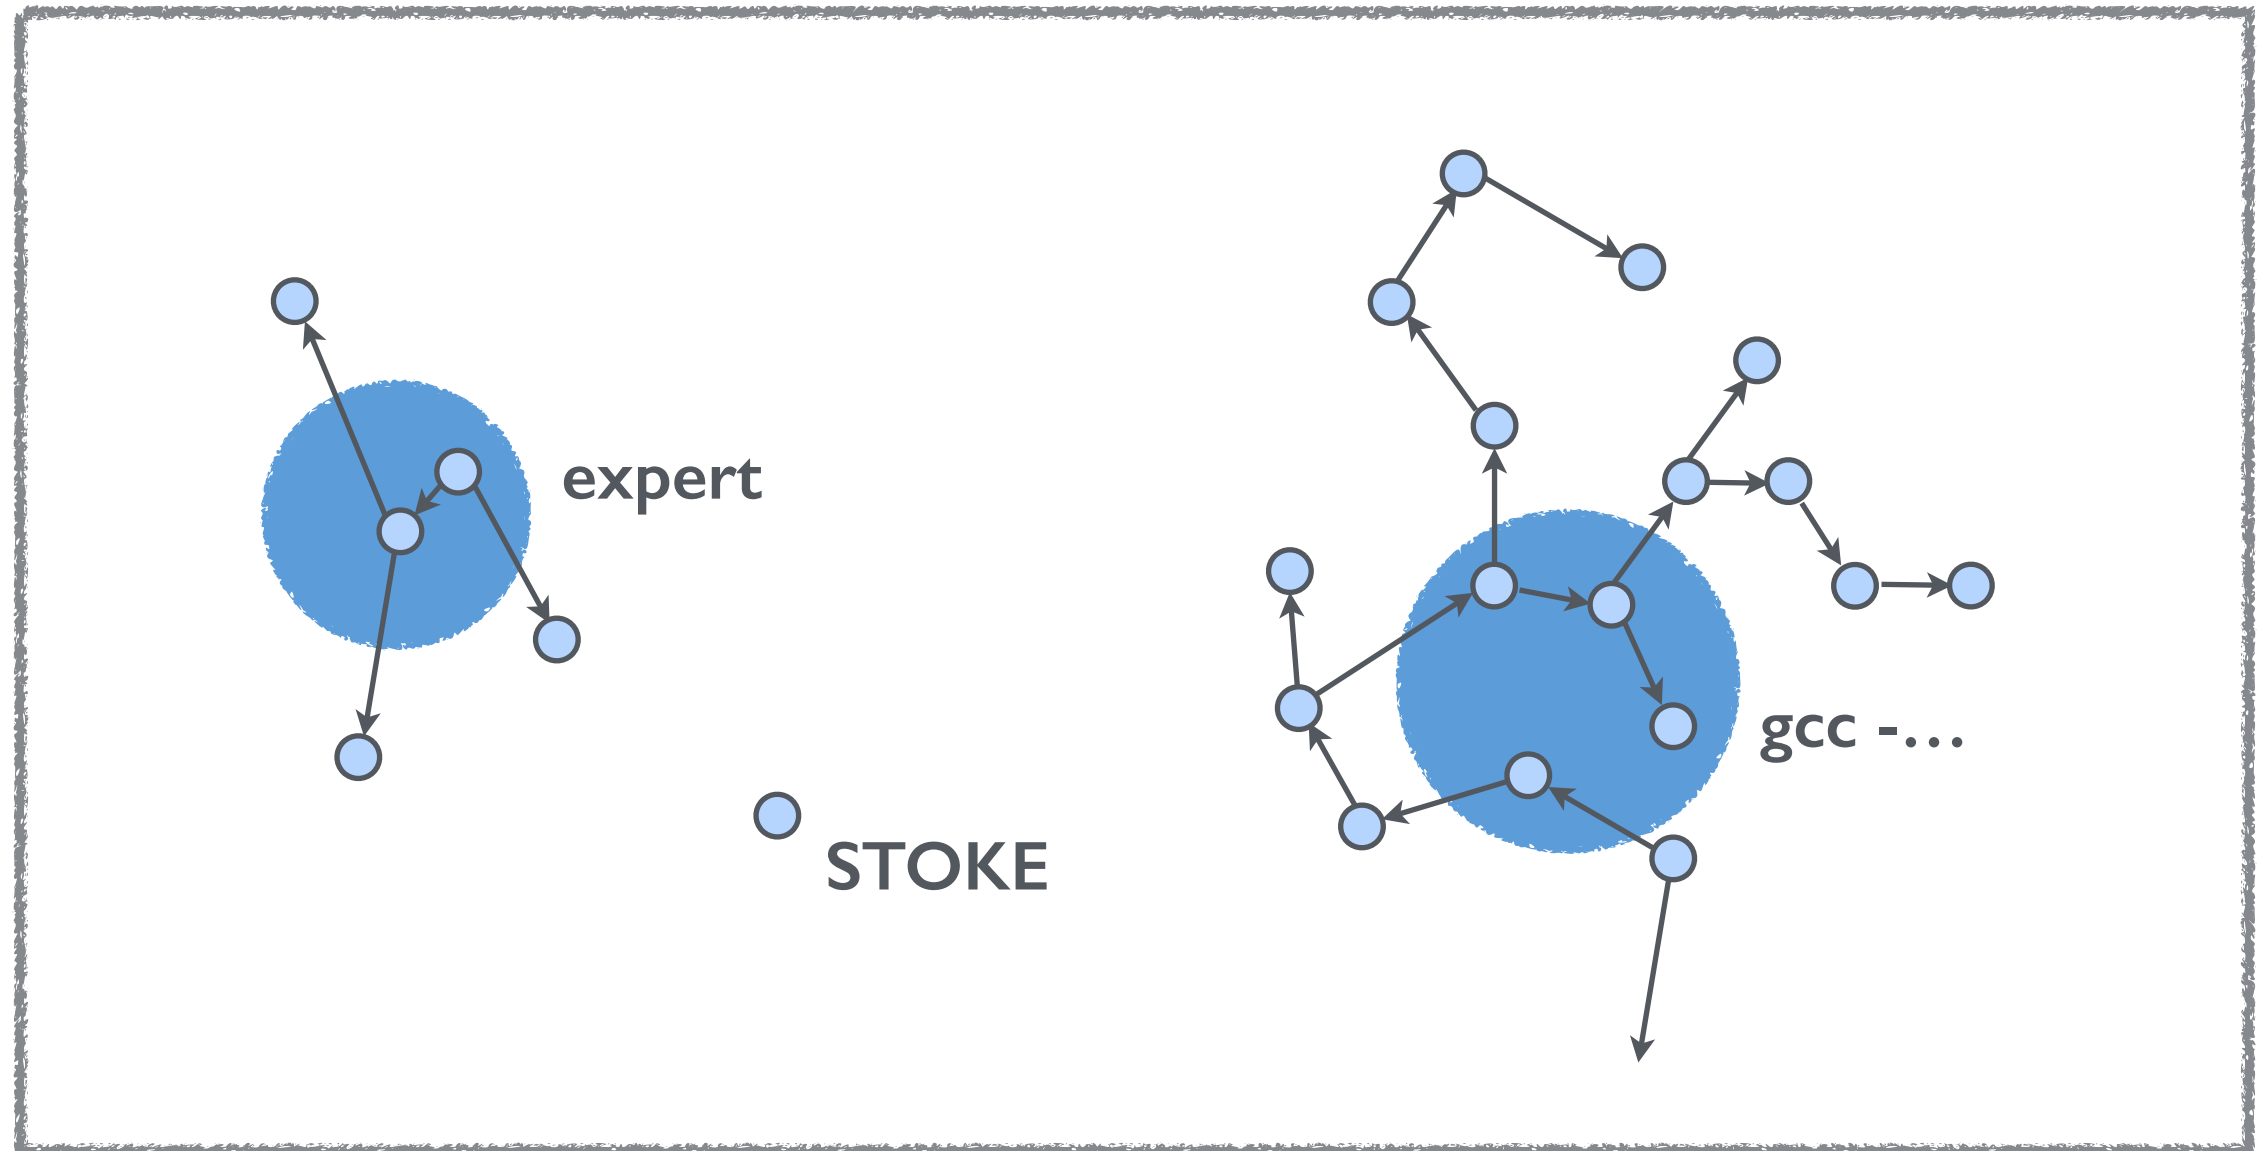
- **Smaller kernel:** 38 LOC reduced to 28 LOC
- **Performance improvement:** 57% kernel speedup, produces a 27% overall task speedup
- **Highly specialized:** Obeys application-specific error bound requirements for all inputs between -3.0 and 0

# Intuition



**Abstract space of programs (*rewrites*):** Blue regions contains points that are bit-wise equivalent to the *target*

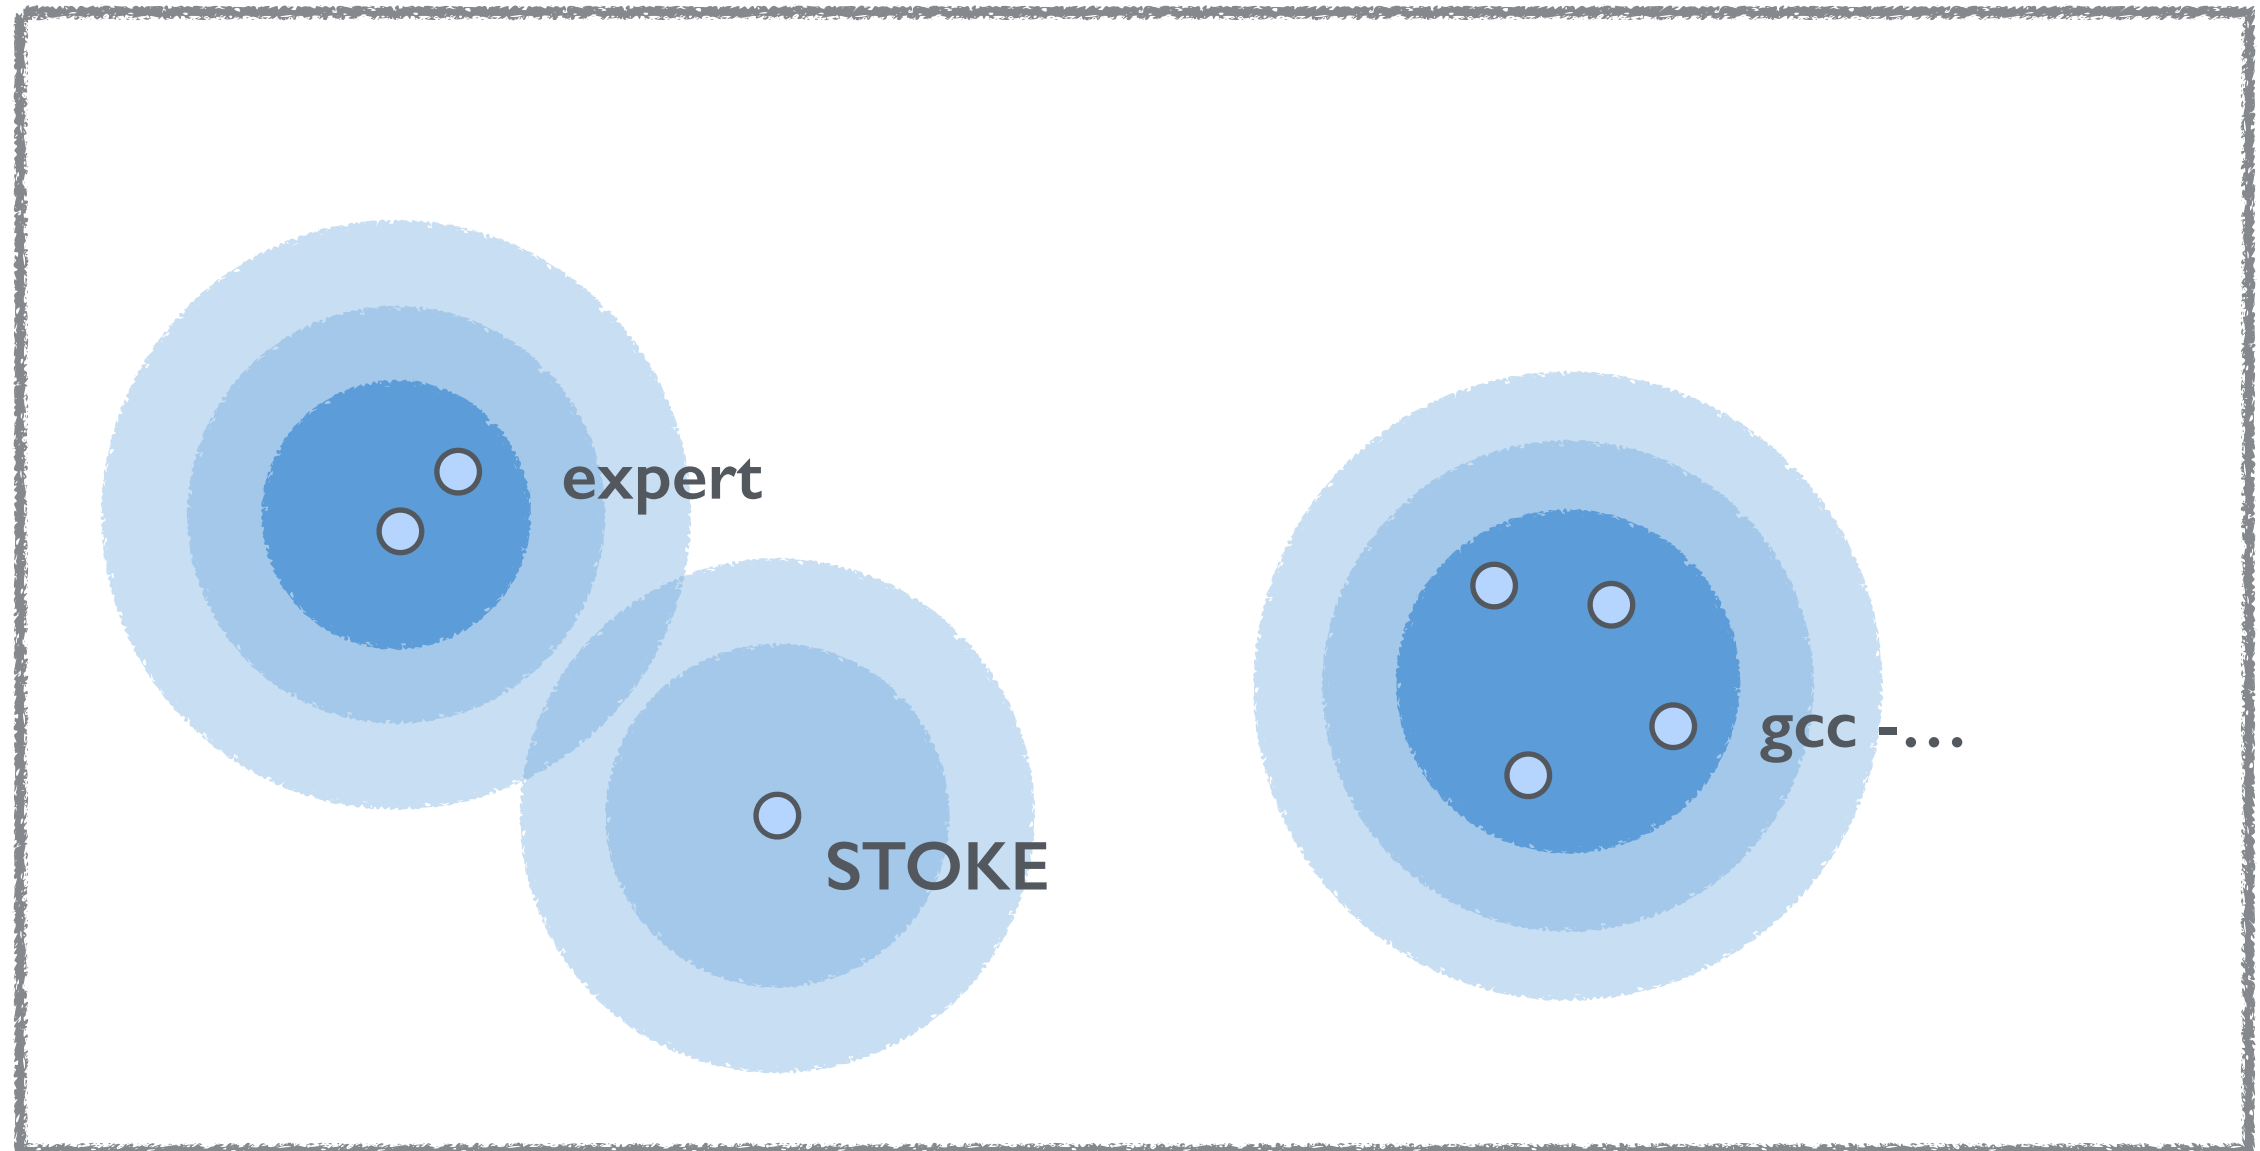
# Intuition



**Few opportunities:** Floating-point instruction set semantics complicate optimization

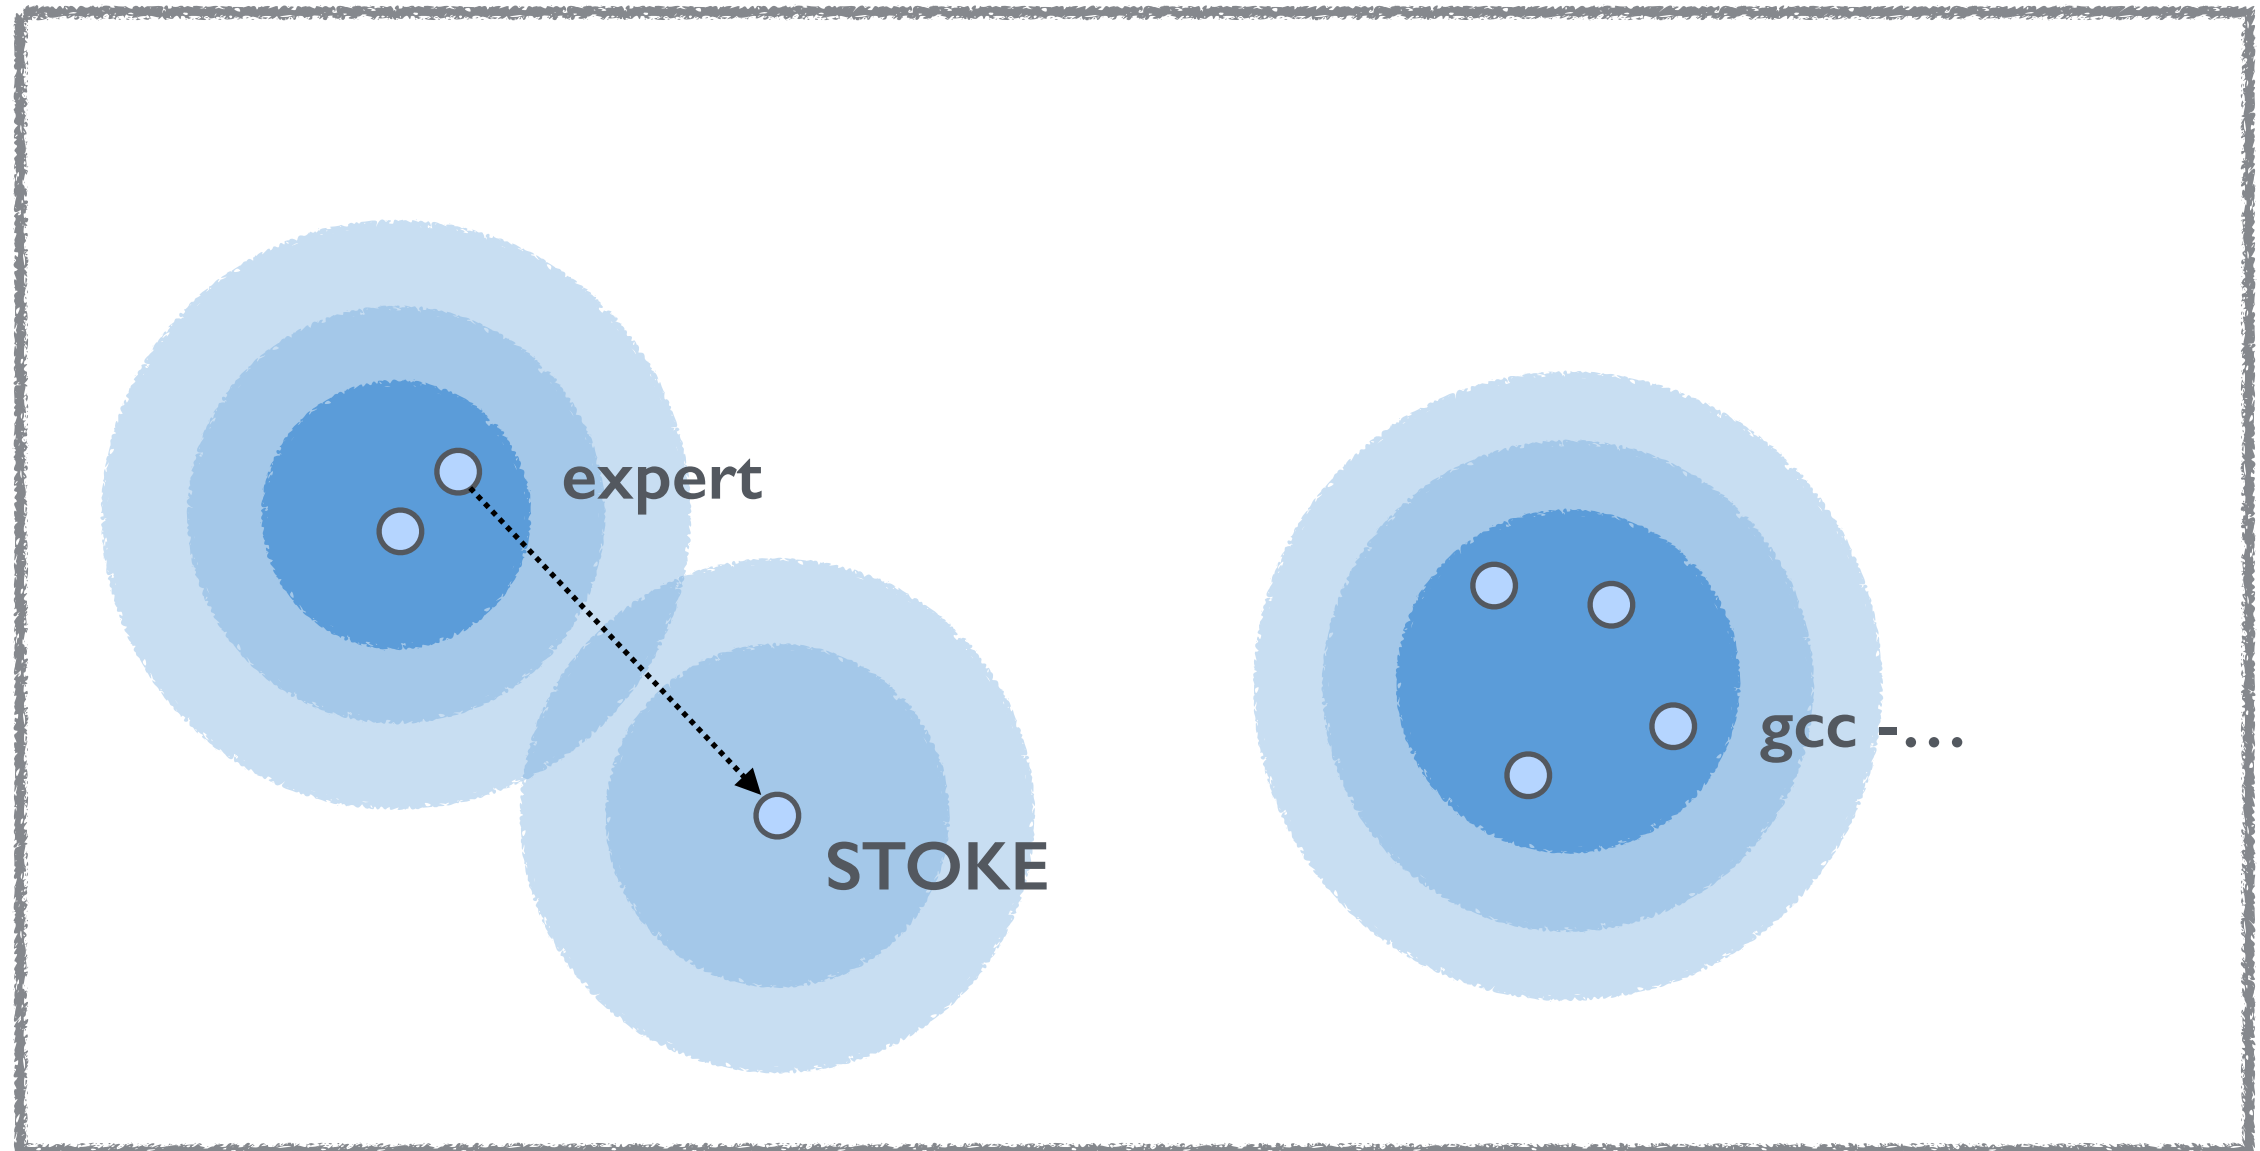


# Intuition



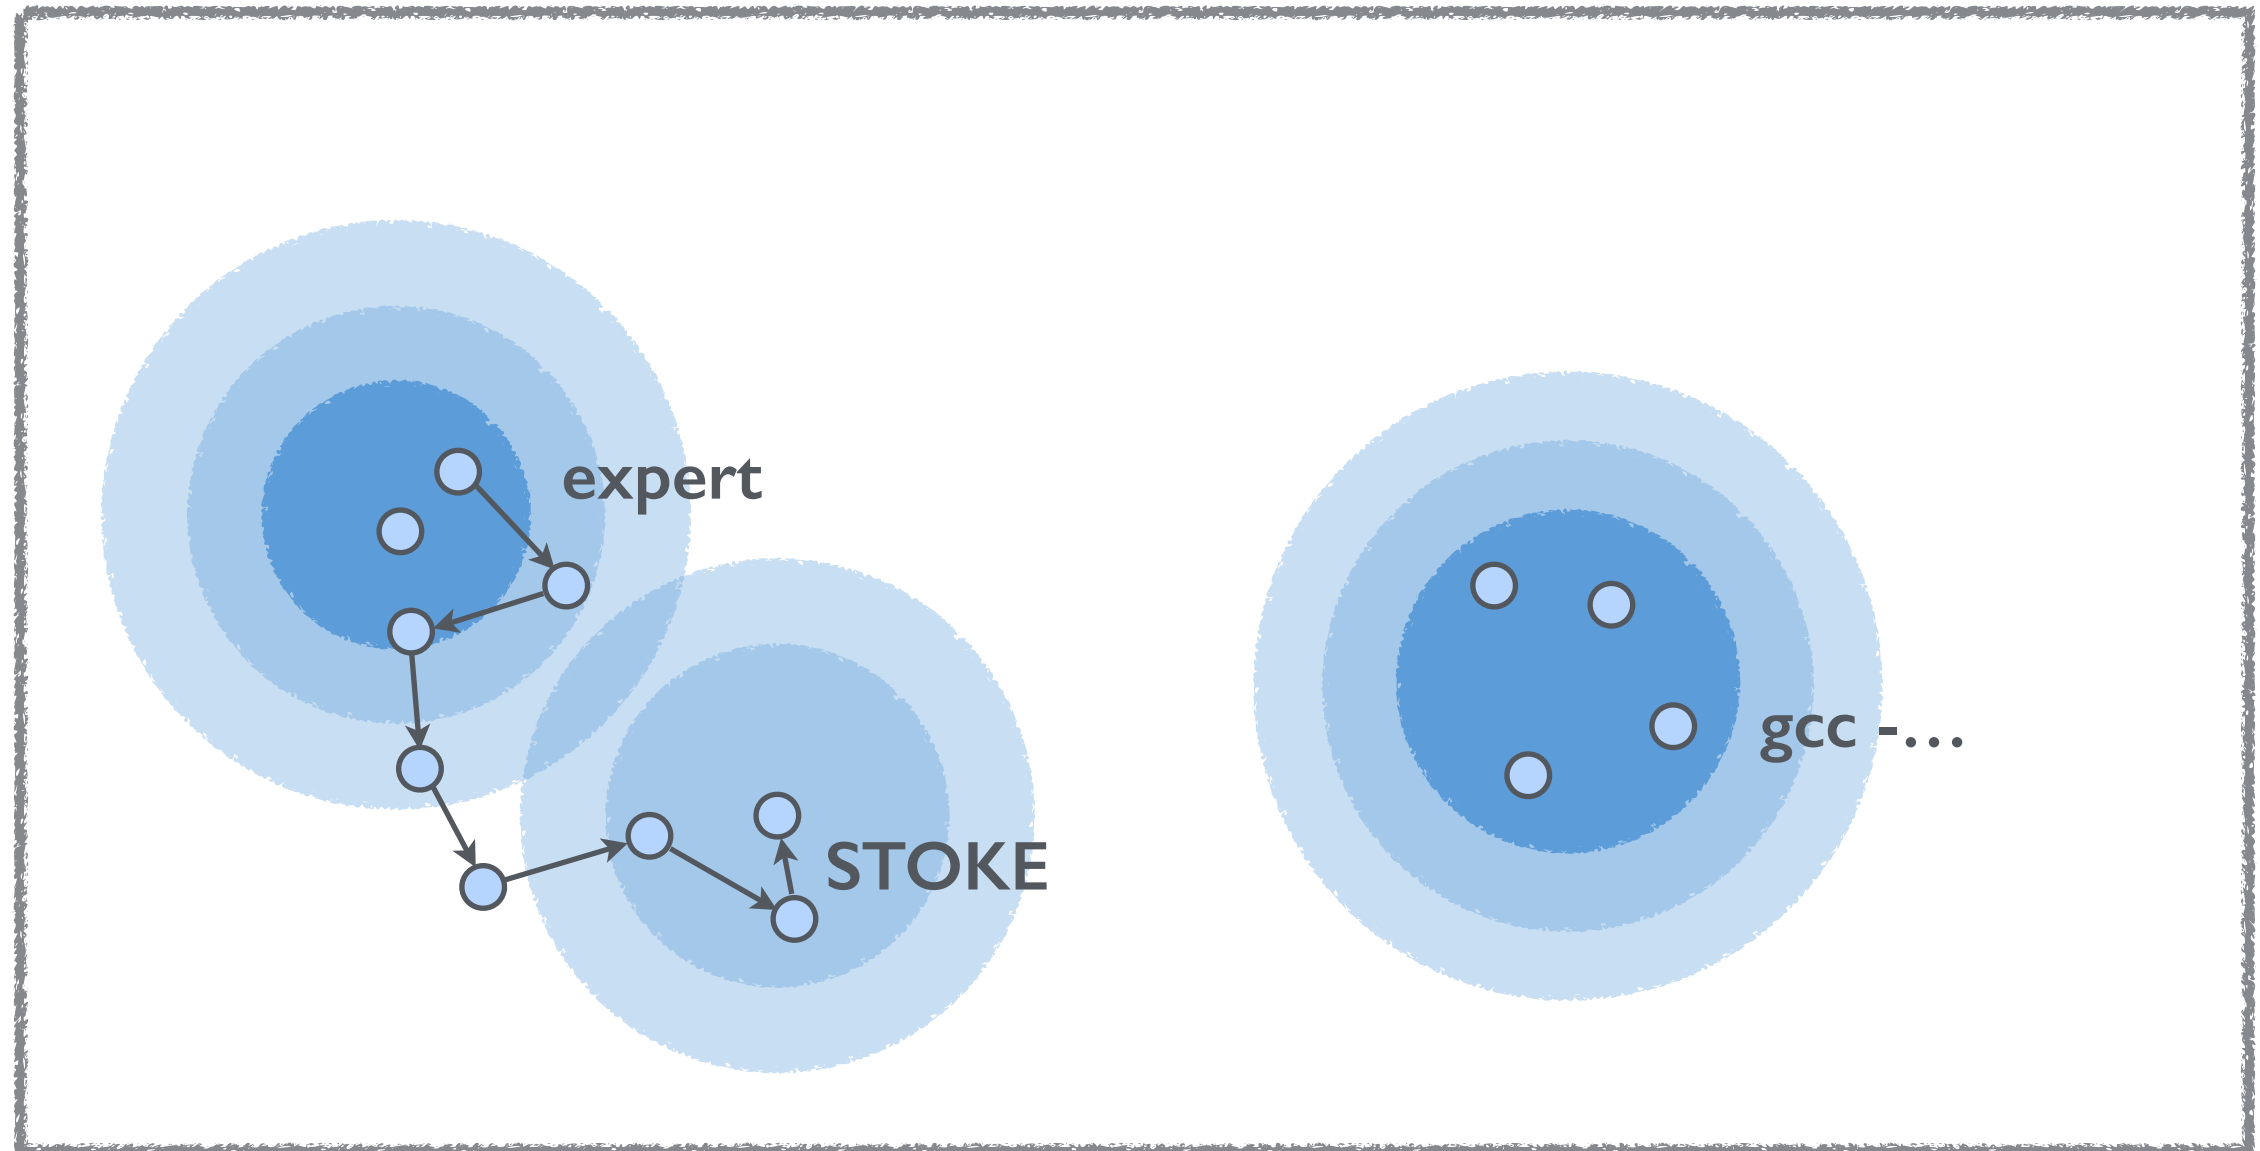
**Relax correctness:** Most applications don't require full precision results for all possible inputs

# Intuition



**New opportunities:** Gain access to high performance optimizations that were previous inaccessible

# Intuition



**Non-trivial task:** Semantics preserving transformations ill-suited to this task; prefer stochastic search [schkufza 13]

# What's Required

1. **Search Procedure:** Markov Chain Monte Carlo Sampling
2. **Cost Function:** Formal encoding of rewrite quality to guide search; should balance competing constraints of precision and performance

# MCMC Sampling

- **Widely used:** For many domains, the only known tractable solution method for high dimensional irregular search spaces [andrieu 03][chenney 00]
- **Guarantees:** Draws samples in proportion to their value; higher value points are sampled more frequently
- **No claim of convergence rate:** Works well in practice for the benchmarks that we consider

# MCMC Sampling

## Algorithm:

1. Select an initial program
2. Repeat (**millions** to **billions** of times)
  - A. Propose a random change and evaluate cost
  - B. If ( decreased ) { accept }
  - C. If ( increased ) { maybe accept anyway }

# Transformations

## original

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```

# Transformations

## insert

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
imulq rsi, rdx  
...
```



## original

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```



# Transformations

## insert

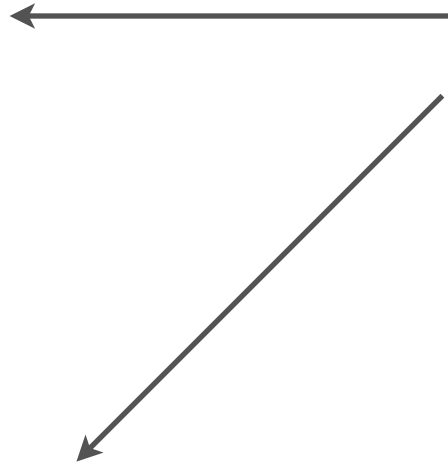
```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
imulq rsi, rdx  
...
```

## delete

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```

## original

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```



# Transformations

## insert

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
imulq rsi, rdx  
...
```

## delete

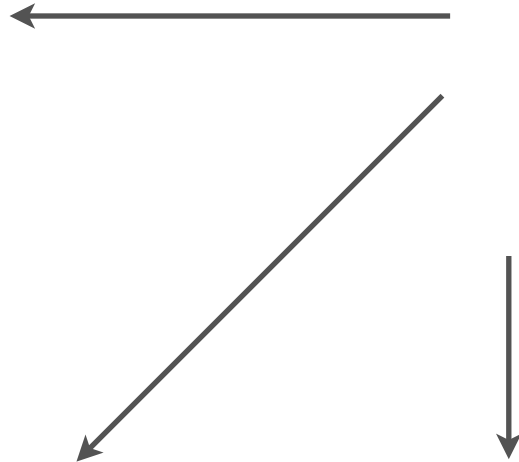
```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```

## original

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```

## instruction

```
...  
movl ecx, ecx  
shrq 32, rsi  
salq 16, rcx  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```



# Transformations

## insert

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
imulq rsi, rdx  
...
```

## delete

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```

## original

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```

## instruction

```
...  
movl ecx, ecx  
shrq 32, rsi  
salq 16, rcx  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```

## opcode

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
subl edx, edx  
imulq r9, rax  
...
```

# Transformations

## insert

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
imulq rsi, rdx  
...
```

## delete

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```

## original

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```

## instruction

```
...  
movl ecx, ecx  
shrq 32, rsi  
salq 16, rcx  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```

## opcode

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
subl edx, edx  
imulq r9, rax  
...
```

## operand

```
...  
movl ecx, ecx  
shrq 32, rcx  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```

# Transformations

## insert

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
imulq rsi, rdx  
...
```

## delete

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```

## original

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```

## instruction

```
...  
movl ecx, ecx  
shrq 32, rsi  
salq 16, rcx  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```

## swap

```
...  
movl ecx, ecx  
movl edx, edx  
andl ff, r9d  
movq rcx, rax  
shrq 32, rsi  
imulq r9, rax  
...
```

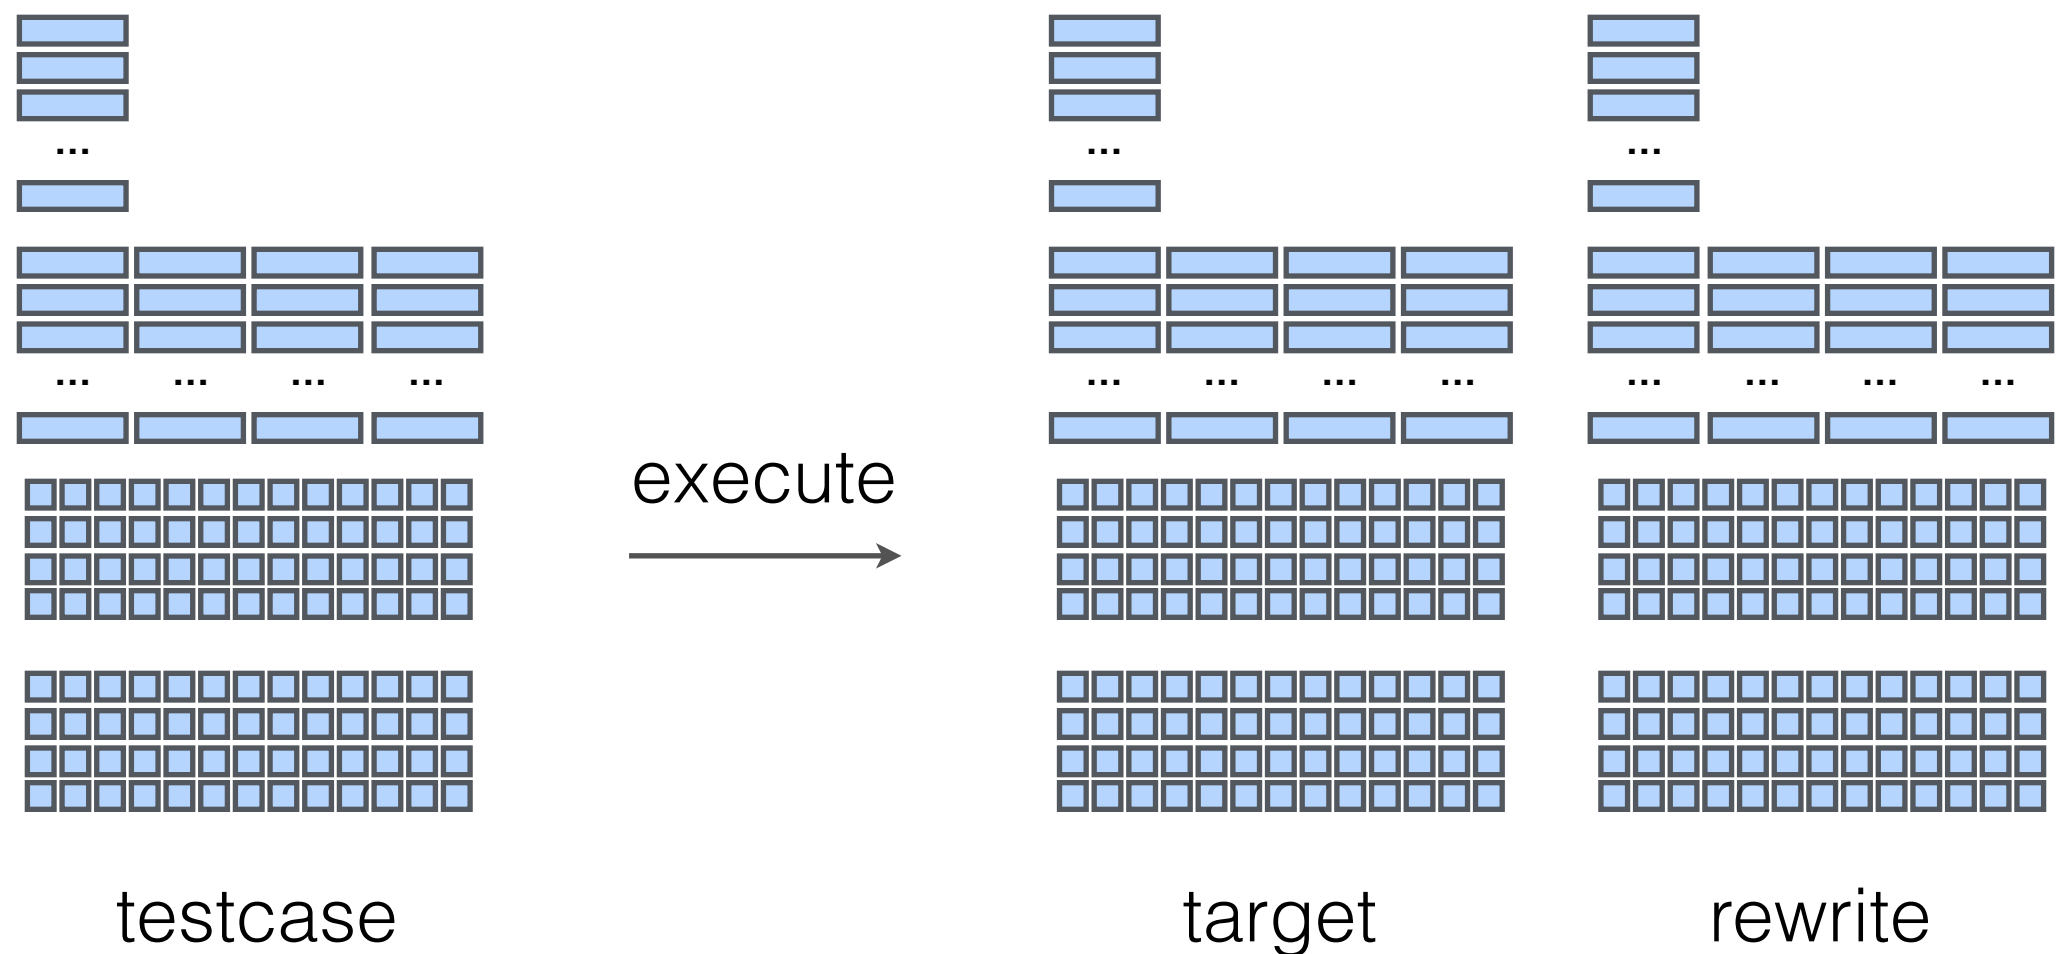
## opcode

```
...  
movl ecx, ecx  
shrq 32, rsi  
andl ff, r9d  
movq rcx, rax  
subl edx, edx  
imulq r9, rax  
...
```

## operand

```
...  
movl ecx, ecx  
shrq 32, rcx  
andl ff, r9d  
movq rcx, rax  
movl edx, edx  
imulq r9, rax  
...
```

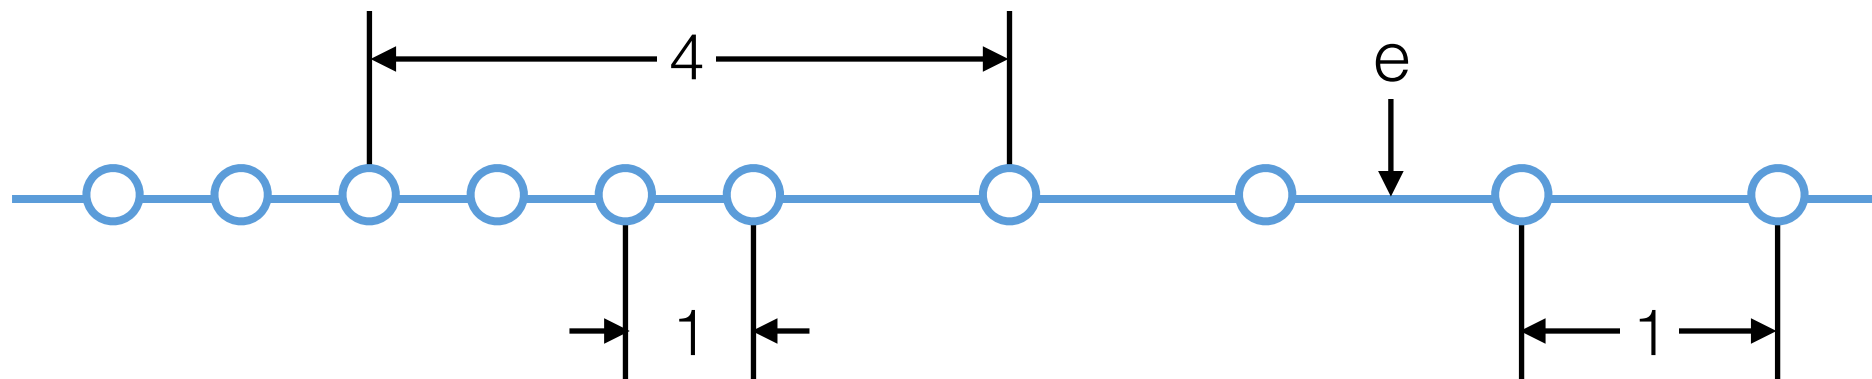
# Precision Function



**Comparison:** Execute target and rewrite on identical state and compare live outputs

# Precision Function

- **Uncertainty in Last Place:** Measures the distance between a real number and the closest representable floating-point value
- **Widely Used:** Most scientific applications measure precision in terms of ULPs; 0.5 is the gold standard but very expensive to obtain, most settle for 1 to 2



# Verification

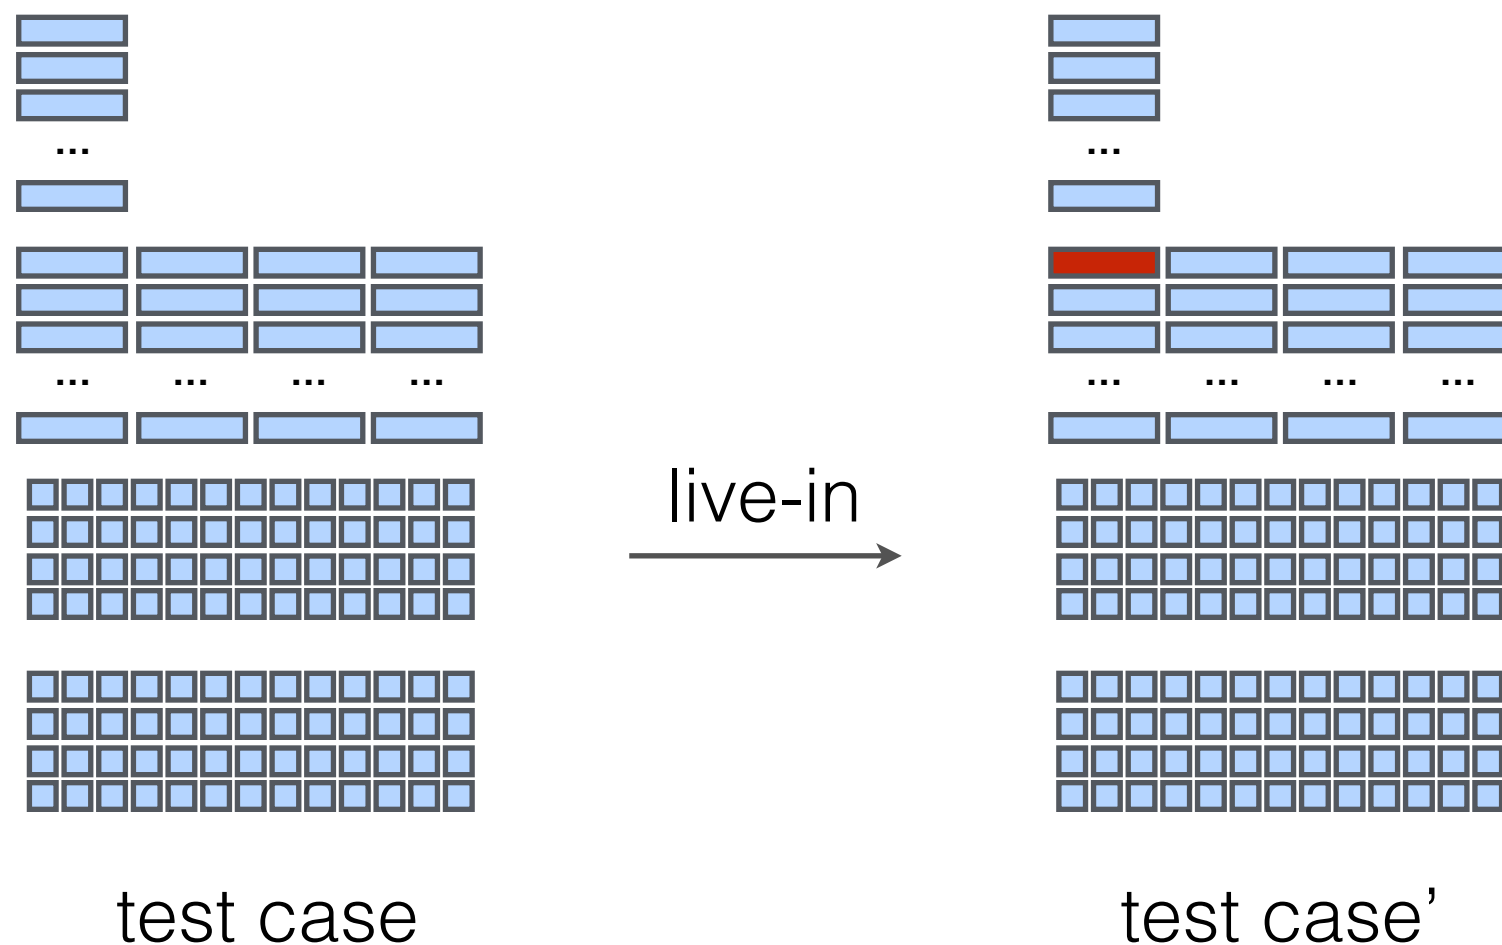
- **Guarantees:** How do we know that an optimization is “precise enough” for all possible inputs?
- **Decision Procedures:** Bit-blasting techniques don’t scale beyond a few lines of code; can’t handle mixed fixed- and floating-point codes [darulova 14]
- **Abstract Interpretation:** Can’t prove even bitwise equality in many common cases; can’t handle mixed fixed- and floating-point codes [haller 12]
- **Bottom Line:** No standard techniques can do this



# Validation

- **Relaxed problem statement:** Claim with high confidence that there is no input that will cause an error in excess of maximum user bound
- **Error function:** Run original and optimized code on identical inputs and measure ULP Error
$$\text{error}(x) = \text{ULP}(\text{eval}(\text{target}, x), \text{eval}(\text{rewrite}, x))$$
- **Goal:** Show max error is below user-defined bound

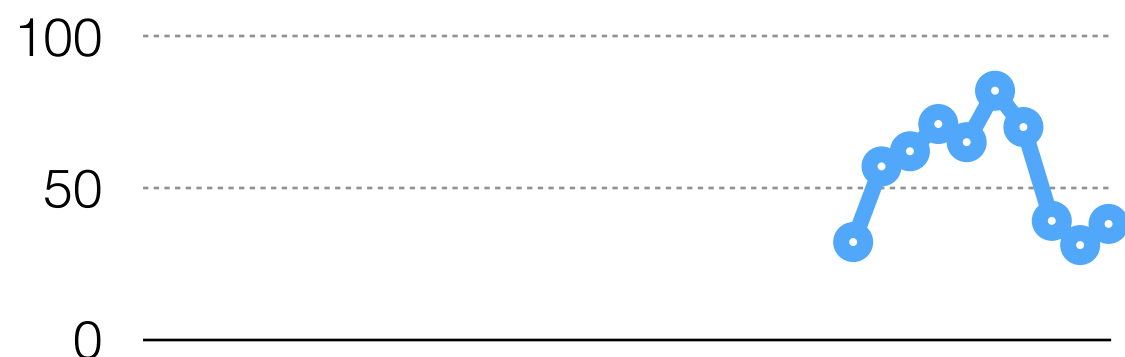
# Validation



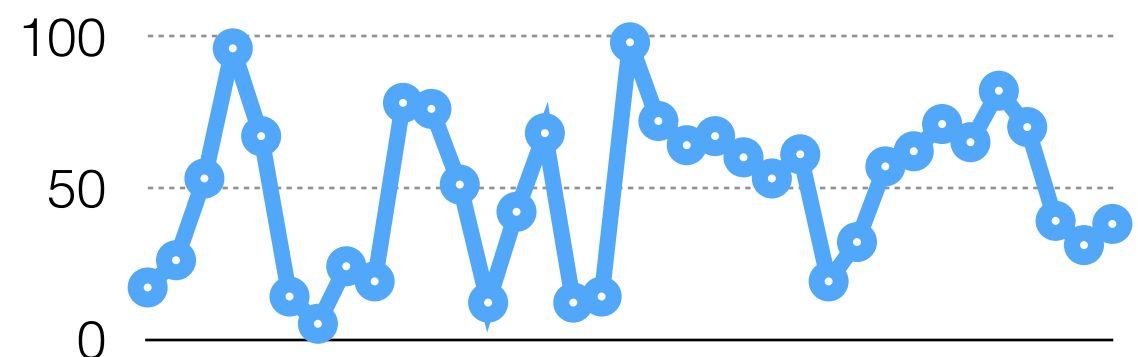
**Search:** Use MCMC sampling to search for test case inputs that maximize the error function; just one transform

# Termination

- **Mixing Tests:** Statistical tests [geweke 92] for producing a high-confidence guarantee that search has sampled uniformly across the domain of a function
- **Therefore:** High confidence that search has discovered all local maxima implies high confidence that it has discovered the global maximum



not converged

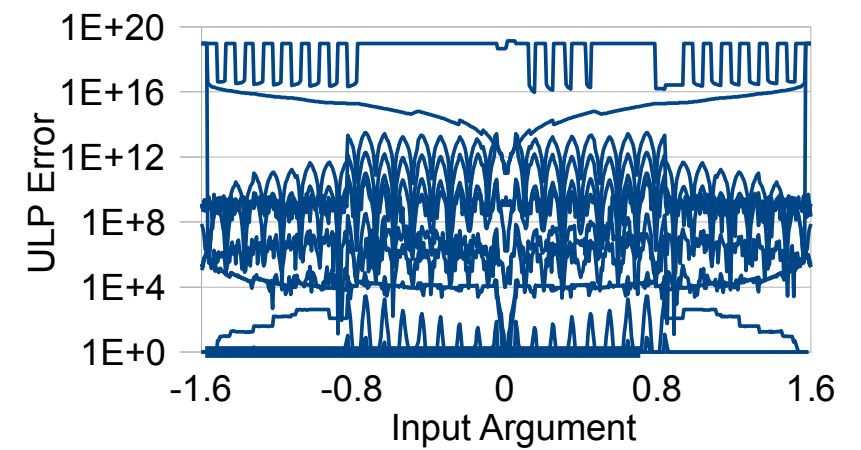
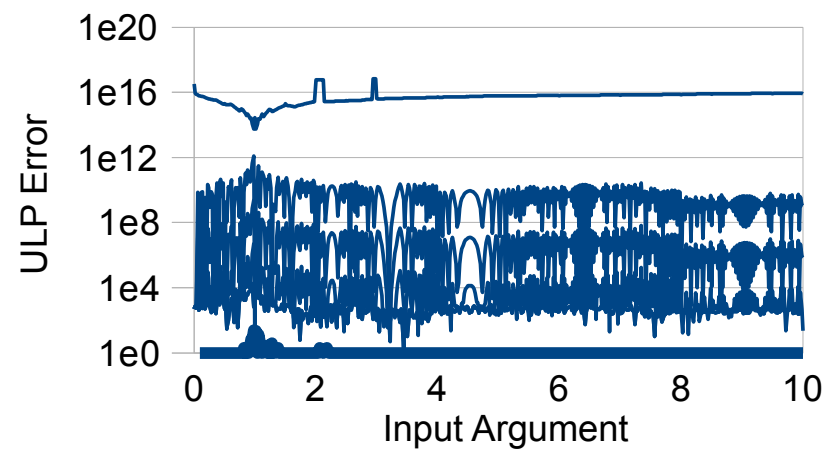
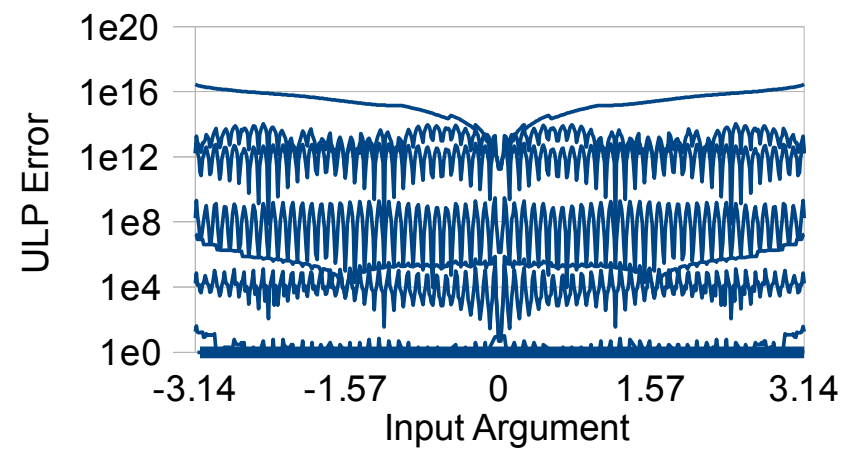
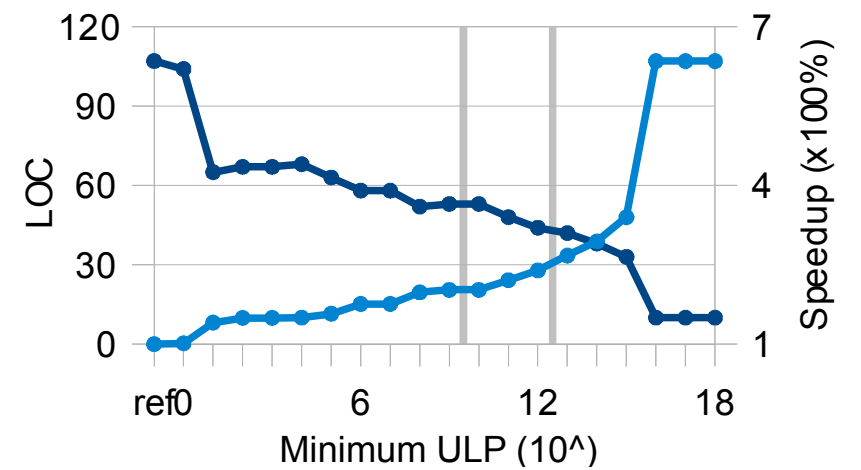
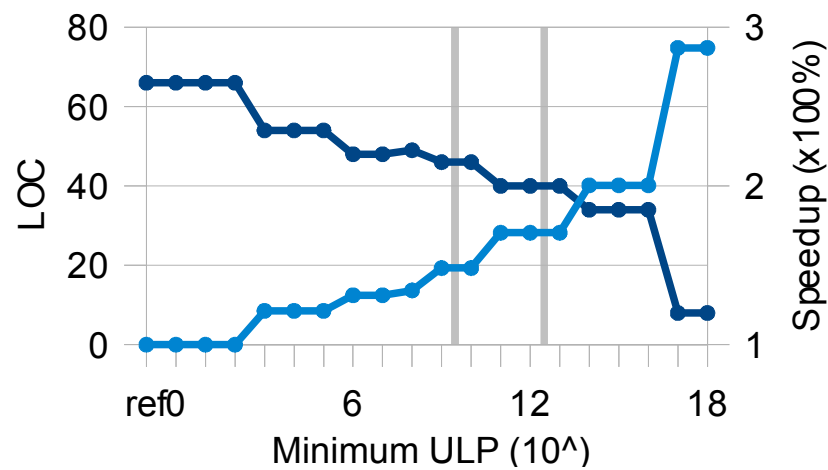
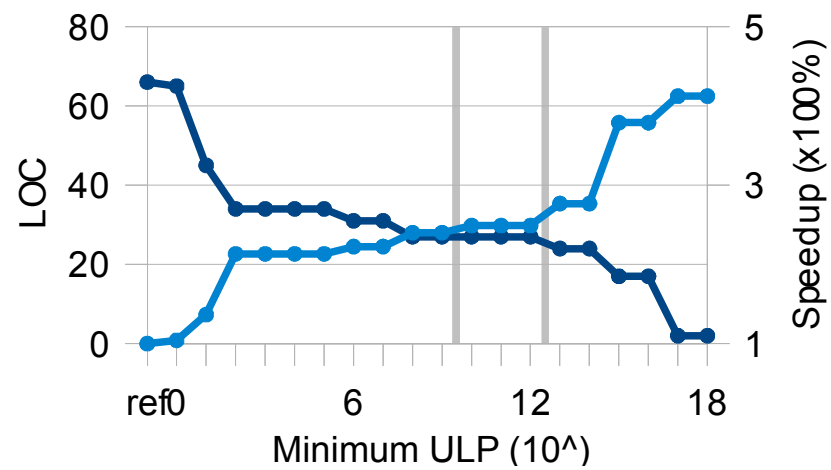


converged

# Evaluation

- **Numeric simulation:** S3D, 3-dimensional direct numerical solver for HCCI combustion
- **C library:** libimf, Intel's hand-written implementation of the C numerics library math.h
- **Computer graphics:** A ray tracer

# libimf



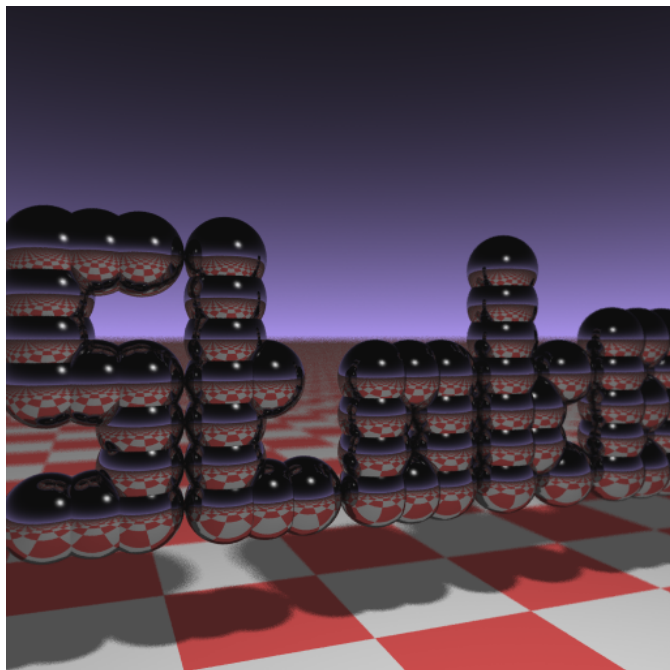
$\sin()$

$\log()$

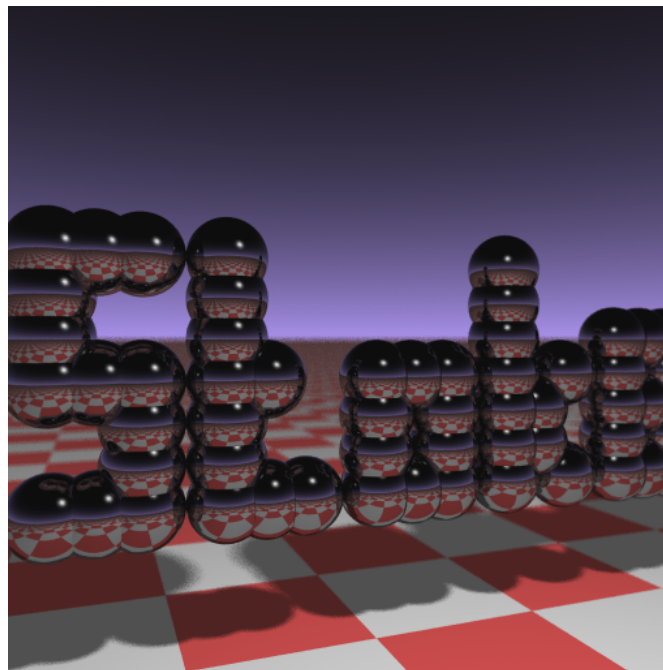
$\tan()$

# Ray Tracer

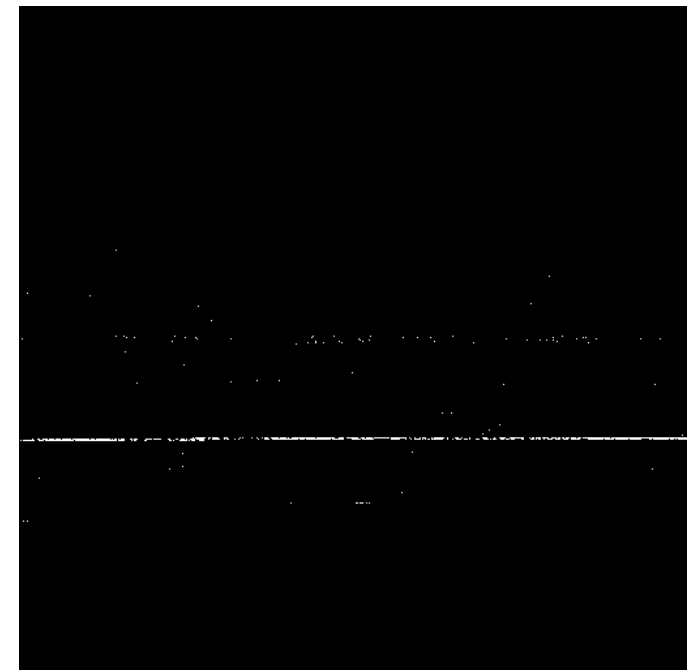
- **Bit-wise correct:** 30% speedup optimizing vector kernels
- **Depth of field blur:** Random perturbations made to viewing camera angle; 6% speedup by relaxing precision requirements



bit-wise correct



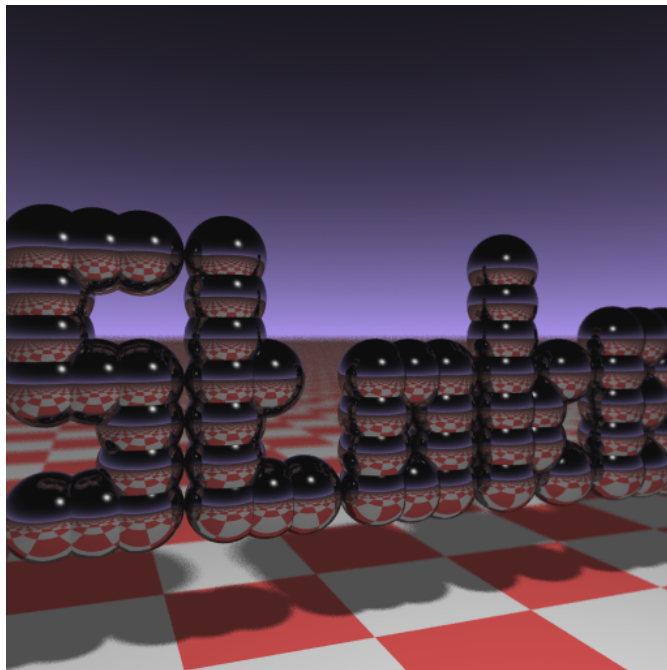
relaxed precision



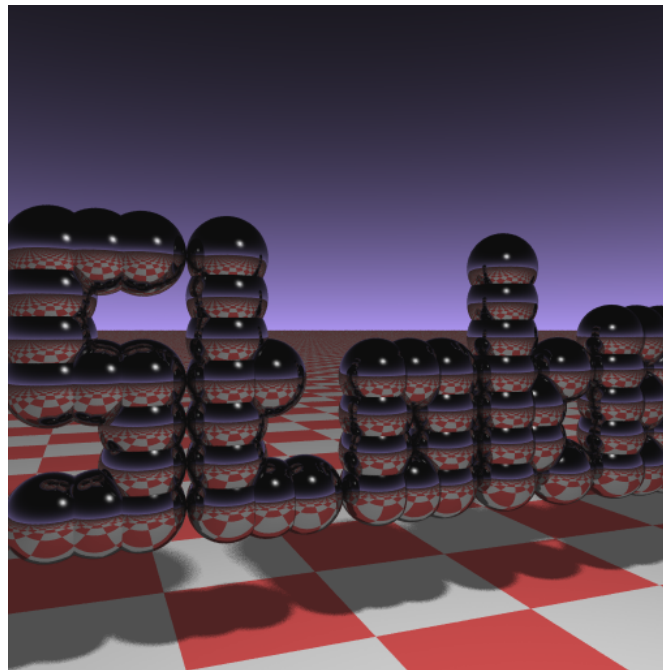
error pixels (white)

# Overfitting

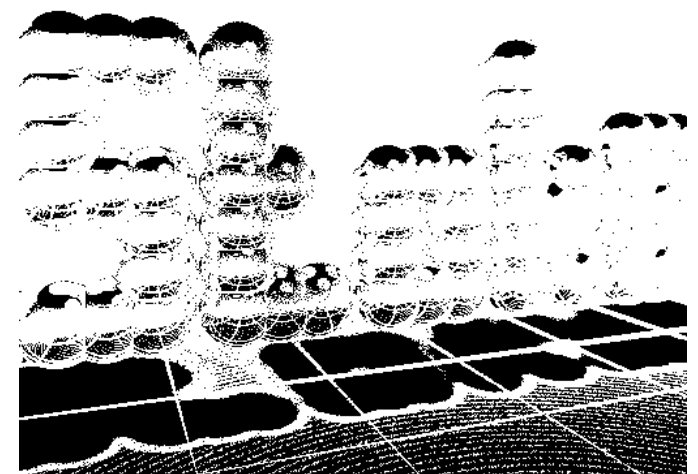
- **Over-relaxation:** If minimum tolerable error exceeds the variance of random perturbations, they are removed altogether
- **Faster still:** But depth of field blur has been removed



bit-wise correct



over-relaxed precision



error pixels (white)

# Summary

- **Micro-optimization:** For many interesting application domains, once data movement is orchestrated correctly, even a single instruction can make a difference
- **New approach:** Use random search to experiment with imprecise intermediate optimizations
- **New opportunities for optimization:** Identify applications that can tolerate a loss of precision and produce code that is specialized to that domain
- **Download:** [www.github.com/eschkufz/stoke-release](https://www.github.com/eschkufz/stoke-release)